

CA1: Dataframe Manipulation with Spotify Data

Jony Karmakar

Introduction

Pandas is an extremely powerful tool to handle large amounts of tabular data. In this compulsory assignment, you will use Pandas to explore one of the TA's personal spotify data in depth. Additional information:

- Feel free to create additional code cells if you feel that one cell per subtask is not sufficient.
- Remember, Pandas uses very efficient code to handle large amounts of data. For-loops are not efficient. If you ever have to use a for-loop to loop over the rows in the DataFrame, you have *probably* done something wrong.
- Label all graphs and charts if applicable.

Task

I typically enjoy indie and rock music. I am a big fan of everything from old-fashioned rock and roll like Led Zeppelin and Jimi Hendrix, to newer indie artists like Joji and Lana Del Rey. This is why my spotify wrapped for 2023 came as quite a surprise:

Now, I'm no hater of pop music, but this was unexpected. For this assignment, you will investigate my listening habits, including a deep dive into my Ariana Grande listening habits, and try to find an answer to why she was my top artist; was there a fault in the spotify algorithm? Am I actually secretly an *Arianator*? (yes, I did have to look that up). Or am I just lying to myself about how often I listen to guilty pleasure music?

Part 1: Initial loading and exploration

1.0 Import necessary libraries:

pandas, numpy, matplotlib.pyplot (other libraries such as seaborn or plotly are also allowed if you want prettier plots). It might also be a good idea to use **os** for task 2.0

```
import matplotlib.pyplot as plt
# ---- Insert other imports ----
import pandas as pd
import numpy as np
import seaborn as sns
```

1.1 Loading the data

Load the dataset in the file `streaming_history_0.csv` into a Pandas DataFrame called `df_spotify_0`.

```
# Loading the dataset
```

```
df_spotify_0 =  
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghistory0.csv')
```

1.2 Help function

Use the Python command `help` to help you understand how to use the `pd.DataFrame.head` and `pd.DataFrame.tail` methods.

```
# Using help command
```

```
help(pd.DataFrame.head)
```

Help on function head in module pandas.core.generic:

```
head(self, n: 'int' = 5) -> 'Self'  
    Return the first `n` rows.
```

This function returns the first `n` rows for the object based on position. It is useful for quickly testing if your object has the right type of data in it.

For negative values of `n`, this function returns all rows except the last `|n|` rows, equivalent to `df[:n]`.

If `n` is larger than the number of rows, this function returns all rows.

Parameters

`n` : int, default 5
 Number of rows to select.

Returns

same type as caller
 The first `n` rows of the caller object.

See Also

`DataFrame.tail`: Returns the last `n` rows.

Examples

```
>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon',  
'lion',  
...                               'monkey', 'parrot', 'shark', 'whale',  
'zebra']})  
>>> df
```

	animal
0	alligator
1	bee
2	falcon
3	lion
4	monkey
5	parrot
6	shark
7	whale
8	zebra

Viewing the first 5 lines

```
>>> df.head()
```

	animal
0	alligator
1	bee
2	falcon
3	lion
4	monkey

Viewing the first `n` lines (three in this case)

```
>>> df.head(3)
```

	animal
0	alligator
1	bee
2	falcon

For negative values of `n`

```
>>> df.head(-3)
```

	animal
0	alligator
1	bee
2	falcon
3	lion
4	monkey
5	parrot

Using help command
`help(pd.DataFrame.tail)`

Help on function tail in module pandas.core.generic:

```
tail(self, n: 'int' = 5) -> 'Self'  
    Return the last `n` rows.
```

This function returns last `n` rows from the object based on position. It is useful for quickly verifying data, for example, after sorting or appending rows.

For negative values of `n`, this function returns all rows except the first `|n|` rows, equivalent to ``df[|n|:]``.

If n is larger than the number of rows, this function returns all rows.

Parameters

n : int, default 5
 Number of rows to select.

Returns

type of caller
 The last `n` rows of the caller object.

See Also

DataFrame.head : The first `n` rows of the caller object.

Examples

```
-----  
>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon',  
'lion',  
    ..., 'monkey', 'parrot', 'shark', 'whale',  
'zebra']})  
>>> df  
   animal  
0  alligator  
1      bee  
2   falcon  
3     lion  
4   monkey  
5   parrot  
6   shark  
7   whale  
8   zebra
```

Viewing the last 5 lines

```
>>> df.tail()
      animal
4  monkey
5  parrot
6   shark
7   whale
8   zebra
```

Viewing the last `n` lines (three in this case)

```
>>> df.tail(3)
      animal
6   shark
7   whale
8   zebra
```

For negative values of `n`

```
>>> df.tail(-3)
      animal
3    lion
4  monkey
5  parrot
6   shark
7   whale
8   zebra
```

1.3 Getting an overview

Print the first `five` and last `ten` rows of the dataframe. Have a quick look at which columns are in the dataset.

```
print(df_spotify_0.head())
print(df_spotify_0.tail(10))
```

	endTime	artistName
trackName \		
0 2022-12-03 02:02	Cigarettes After Sex	
Truly		
1 2022-12-03 02:02	Leonard Cohen	Take This Waltz - Paris
Version		
2 2022-12-06 21:05	Vlad Holiday	So Damn Into
You		
3 2022-12-06 21:05	Lorde	
Team		
4 2022-12-06 21:05	Ariana Grande	Into
You		

```

    msPlayed
0    30000.0
1     8210.0
2    37895.0
3     8984.0
4     1221.0

    endTime    artistName
trackName \
11949  2023-01-02 20:58    Ariana Grande          six
thirty
11950  2023-01-02 20:58    Leonard Cohen      Thanks for the
Dance
11951  2023-01-02 20:59          Des Rocs      Used to the
Darkness
11952  2023-01-02 20:59  Caroline Polachek      Hit Me Where It
Hurts
11953  2023-01-02 20:59  Caroline Polachek      Hit Me Where It
Hurts
11954  2023-01-02 20:59  Kaizers Orchestra
Resistansen
11955  2023-01-02 20:59          Mr.Kitty          After
Dark
11956  2023-01-02 20:59    daddy's girl  after dark x sweater
weather
11957  2023-01-02 20:59    daddy's girl  after dark x sweater
weather
11958  2023-01-02 20:59    daddy's girl  after dark x sweater
weather

    msPlayed
11949    1699.0
11950   19483.0
11951     185.0
11952     603.0
11953     208.0
11954     208.0
11955  101447.0
11956     301.0
11957     208.0
11958     789.0

```

1.4 Formatting correctly

When working with Pandas, it's very useful to have columns which contains dates in a specific format called *datetime*. This allows for efficient manipulation and analysis of time-series data, such as sorting, filtering by date or time, and resampling for different time periods. Figure out which column(s) would be appropriate to convert to datetime, if any, and if so, perform the conversion to the correct format.

```
# Formatting the endTime column to datetime dtype of pandas
df_spotify_0['endTime'] = pd.to_datetime(df_spotify_0['endTime'])

df_spotify_0.head()
```

	trackName \	endTime	artistName
0	2022-12-03 02:02:00	Cigarettes After Sex Truly	
1	2022-12-03 02:02:00	Leonard Cohen	Take This Waltz - Paris Version
2	2022-12-06 21:05:00	Vlad Holiday	So Damn Into You
3	2022-12-06 21:05:00	Lorde	Team
4	2022-12-06 21:05:00	Ariana Grande	Into You

	msPlayed
0	30000.0
1	8210.0
2	37895.0
3	8984.0
4	1221.0

1.5 Unique artists

Find how many unique artists are in the dataset.

```
# Finding and showing the unique artists number
artist_name = df_spotify_0['artistName'].unique()
artist_name
print(f'Total No. Unique Artist: {len(artist_name)}')
```

Total No. Unique Artist: 495

1.6 Unique songs

Find how many unique songs are in the dataset.

```
# Funding and showing the unique artists number
songs_name = df_spotify_0['trackName'].unique()
songs_name
print(f'Total No. Unique Songs: {len(songs_name)}')
```

len(songs_name)

Total No. Unique Songs: 1308

1308

Part 1: Questions

Q1: Which columns are in the dataset? A1: There are 4 columns in the dataset. a. endTime b. artistName c. trackName d. msPlayed Q2: What timeframe does the dataset span? A2: 2022-12-03 02:02 to 2023-01-02 20:59. Q3: How many unique artists are in the dataset? A3: 495 Q4: How many unique songs are in the dataset? A4: 1308

Part 2: Working with all the data

2.0 Importing all the dataframes

In Task 1, you only worked with about a month worth of data. Now, you will work with over a year worth.

In the *spotify_data* folder, there is more than just one listening record. Load each of the 14 listening records into a dataframe (1 dataframe per listening record), and concatenate them together into one large dataframe named *df*.

```
# Loading rest of the dataset
df_spotify_1 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry1.csv')
df_spotify_2 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry2.csv')
df_spotify_3 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry3.csv')
df_spotify_4 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry4.csv')
df_spotify_5 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry5.csv')
df_spotify_6 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry6.csv')
df_spotify_7 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry7.csv')
df_spotify_8 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry8.csv')
df_spotify_9 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry9.csv')
df_spotify_10 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry10.csv')
```



```

df_spotify_11 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry11.csv')
df_spotify_12 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry12.csv')
df_spotify_13 =
pd.read_csv('/Users/jony/DAT200/assignment/spotify_data/streaminghisto
ry13.csv')

```

Formatting the endTime dtype to datetime

```

df_spotify_1['endTime'] = pd.to_datetime(df_spotify_1['endTime'])
df_spotify_2['endTime'] = pd.to_datetime(df_spotify_2['endTime'])
df_spotify_3['endTime'] = pd.to_datetime(df_spotify_3['endTime'])
df_spotify_4['endTime'] = pd.to_datetime(df_spotify_4['endTime'])
df_spotify_5['endTime'] = pd.to_datetime(df_spotify_5['endTime'])
df_spotify_6['endTime'] = pd.to_datetime(df_spotify_6['endTime'])
df_spotify_7['endTime'] = pd.to_datetime(df_spotify_7['endTime'])
df_spotify_8['endTime'] = pd.to_datetime(df_spotify_8['endTime'])
df_spotify_9['endTime'] = pd.to_datetime(df_spotify_9['endTime'])
df_spotify_10['endTime'] = pd.to_datetime(df_spotify_10['endTime'])
df_spotify_11['endTime'] = pd.to_datetime(df_spotify_11['endTime'])
df_spotify_12['endTime'] = pd.to_datetime(df_spotify_12['endTime'])
df_spotify_13['endTime'] = pd.to_datetime(df_spotify_13['endTime'])

```

Making the list of the whole dataframe

```

dataframes = [df_spotify_0, df_spotify_1, df_spotify_2, df_spotify_3,
df_spotify_4, df_spotify_5,
                df_spotify_6, df_spotify_7, df_spotify_8, df_spotify_9,
df_spotify_10,
                df_spotify_11, df_spotify_12, df_spotify_13]

```

Concatenating the all dataframe in df

```

df = pd.concat(dataframes, ignore_index=True)

```

```

df.head(20)

```

		endTime	artistName \
0	2022-12-03	02:02:00	Cigarettes After Sex
1	2022-12-03	02:02:00	Leonard Cohen
2	2022-12-06	21:05:00	Vlad Holiday
3	2022-12-06	21:05:00	Lorde
4	2022-12-06	21:05:00	Ariana Grande
5	2022-12-07	00:21:00	Caroline Polachek
6	2022-12-07	00:21:00	Kaizers Orchestra
7	2022-12-07	00:21:00	Vlad Holiday
8	2022-12-07	00:21:00	Vlad Holiday
9	2022-12-07	00:21:00	Ariana Grande
10	2022-12-07	00:21:00	Pastel Ghost
11	2022-12-07	00:25:00	Caroline Polachek

12	2022-12-07	00:25:00	BØRNS
13	2022-12-07	00:25:00	Vlad Holiday
14	2022-12-07	00:25:00	Kaizers Orchestra
15	2022-12-07	00:25:00	Ariana Grande
16	2022-12-07	00:25:00	BØRNS
17	2022-12-07	00:25:00	Ariana Grande
18	2022-12-07	00:25:00	Royal Deluxe
19	2022-12-07	00:26:00	The Fratellis

	trackName	msPlayed
0	Truly	30000.0
1	Take This Waltz - Paris Version	8210.0
2	So Damn Into You	37895.0
3	Team	8984.0
4	Into You	1221.0
5	Hit Me Where It Hurts	1648.0
6	Proessen	348.0
7	So Damn Into You	1555.0
8	So Damn Into You	1486.0
9	Best Mistake	8824.0
10	Dark Beach	23518.0
11	Hit Me Where It Hurts	1044.0
12	Clouds	3459.0
13	So Damn Into You	170086.0
14	Proessen	1904.0
15	Just A Little Bit Of Your Heart	557.0
16	Seeing Stars	6145.0
17	safety net (feat. Ty Dolla \$ign)	17838.0
18	I'm A Wanted Man	16152.0
19	She's Not Gone Yet But She's Leaving	48715.0

2.1 Sorting by time

Datasets often aren't perfect. One example of an issue that could occur is that the time-based data might not be in chronological order. If this were to happen, the rows in your dataframe could be in the wrong order. To ensure this isn't an issue in your dataframe, you should sort the dataframe in chronological order, from oldest to newest.

```
# Sorting the dataframe based on the endTime column
df = df.sort_values('endTime', ascending=True)
```

2.2 Setting a timeframe

For this investigation, we are only interested in investigating listening patterns from **2023**. Remove any data not from **2023** from the DataFrame.

```
# Taking data of only 2023 to new dataframe
df2 = df[df['endTime'].dt.year == 2023]
```

```
df2.head(20)
```

		endTime	artistName \
10881	2023-01-01	01:17:00	Ariana Grande
10882	2023-01-01	01:17:00	Ariana Grande
10883	2023-01-01	01:17:00	Ariana Grande
10884	2023-01-01	01:17:00	Peach Pit
10885	2023-01-01	01:17:00	Kelly Clarkson
10886	2023-01-01	01:17:00	Kelly Clarkson
10887	2023-01-01	01:17:00	Ariana Grande
10888	2023-01-01	01:17:00	Ariana Grande
10893	2023-01-01	01:19:00	Regina Spektor
10892	2023-01-01	01:19:00	Lorde
10891	2023-01-01	01:19:00	Lorde
10890	2023-01-01	01:19:00	Fools Garden
10889	2023-01-01	01:19:00	The Neighbourhood
10913	2023-01-01	01:21:00	Marvin Gaye
10912	2023-01-01	01:21:00	Cage The Elephant
10911	2023-01-01	01:21:00	Kaizers Orchestra
10910	2023-01-01	01:21:00	MARBLES
10909	2023-01-01	01:21:00	Lorde
10908	2023-01-01	01:21:00	The Smiths
10907	2023-01-01	01:21:00	The Smiths

	trackName	msPlayed
10881	7 rings	139.0
10882	7 rings	487.0
10883	positions	417.0
10884	Being so Normal	2205.0
10885	Santa, Can't You Hear Me	278.0
10886	Santa, Can't You Hear Me	325.0
10887	Santa Baby	12293.0
10888	Right There (feat. Big Sean)	22929.0
10893	Us - 2005 Remaster	16670.0
10892	Solar Power	464.0
10891	Team	371.0
10890	Lemon Tree	301.0
10889	A Moment of Silence	118932.0
10913	Right On	15904.0
10912	That's Right	14745.0
10911	Hjerteknuser	534.0
10910	I'm Not Ready	441.0
10909	Team	905.0
10908	Heaven Knows I'm Miserable Now - 2011 Remaster	1184.0
10907	Heaven Knows I'm Miserable Now - 2011 Remaster	719.0

```
df2.tail(20)
```

	endTime	artistName
trackName \		

167419	2023-12-07 20:49:00	Lana Del Rey	
Radio			
167420	2023-12-07 20:51:00	Mitski	Washing
Machine Heart			
167421	2023-12-07 20:51:00	Pastel Ghost	
Dark Beach			
167422	2023-12-07 20:51:00	Ariana Grande	
needy			
167423	2023-12-07 21:12:00	Tame Impala	
Eventually			
167424	2023-12-07 21:12:00	U2	With Or
Without You			
167425	2023-12-07 21:12:00	grandson	
Darkside			
167432	2023-12-07 21:13:00	Arctic Monkeys	Snap
Out Of It			
167431	2023-12-07 21:13:00	Lana Del Rey	
Ride			
167433	2023-12-07 21:13:00	Childish Gambino	
Les			
167430	2023-12-07 21:13:00	Lana Del Rey	
Ride			
167434	2023-12-07 21:13:00	Lana Del Rey	
Art Deco			
167428	2023-12-07 21:13:00	Ariana Grande	pete
davidson			
167427	2023-12-07 21:13:00	Labrinth	
Formula			
167426	2023-12-07 21:13:00	Lana Del Rey	Young And
Beautiful			
167429	2023-12-07 21:13:00	Arctic Monkeys	Snap
Out Of It			
167435	2023-12-07 21:13:00	Ariana Grande	off the table (with The
Weeknd)			
167437	2023-12-07 21:14:00	Leonard Cohen	Thanks for
the Dance			
167436	2023-12-07 21:14:00	Ariana Grande	
my hair			
167438	2023-12-07 21:17:00	The Vaccines	Your Love Is My
Favourite Band			

	msPlayed
167419	214573.0
167420	128198.0
167421	19752.0
167422	26220.0
167423	318591.0
167424	13626.0
167425	5404.0

167432	615.0
167431	1126.0
167433	2484.0
167430	859.0
167434	38298.0
167428	603.0
167427	882.0
167426	3146.0
167429	1497.0
167435	13448.0
167437	9317.0
167436	23757.0
167438	14661.0

2.3 Deleting rows

Often in Data Science, you will encounter when a row entry has the value *NaN*, indicating missing data. These entries can skew your analysis, leading to inaccurate conclusions. For this task, identify and remove any rows in your DataFrame that contain NaN values. Later in the course, you might encounter other techniques of dealing with missing data, typically referred to as *data imputation*. Here, though, you are just supposed to delete the entire rows with missing data.

```
# Looking is there any null values in the dataframe or not
null_count = df2.isnull().sum().sum()
print('Number of null values:', null_count)

Number of null values: 18

# Removing the row with null values and storing in a different
dataframe
df3 = df2.dropna()

# Again checking is there any null values in the dataframe or not
null_count = df3.isnull().sum().sum()
print('Number of null values:', null_count)

Number of null values: 0
```

2.4 Convert from milliseconds to seconds

From `msPlayed`, create a new column `secPlayed` with the data converted from milliseconds to seconds. Then delete the column `msPlayed`.

```
# Adding new column named secPlayed which is conversion of miliseconds
to seconds

secPlayed = []

for i in df3['msPlayed']:
```

```
i = i/1000
secPlayed.append(i)
```

```
df3.insert(4, 'secPlayed', secPlayed)
```

```
df3.head()
```

		endTime	artistName	trackName
msPlayed \				
10881	2023-01-01 01:17:00	Ariana Grande	7 rings	
139.0				
10882	2023-01-01 01:17:00	Ariana Grande	7 rings	
487.0				
10883	2023-01-01 01:17:00	Ariana Grande	positions	
417.0				
10884	2023-01-01 01:17:00	Peach Pit	Being so Normal	
2205.0				
10885	2023-01-01 01:17:00	Kelly Clarkson	Santa, Can't You Hear Me	
278.0				

	secPlayed
10881	0.139
10882	0.487
10883	0.417
10884	2.205
10885	0.278

```
# Dropping the msPlayed column
```

```
df3 = df3.drop('msPlayed', axis=1)
```

```
df3.head()
```

		endTime	artistName	trackName
secPlayed				
10881	2023-01-01 01:17:00	Ariana Grande	7 rings	
0.139				
10882	2023-01-01 01:17:00	Ariana Grande	7 rings	
0.487				
10883	2023-01-01 01:17:00	Ariana Grande	positions	
0.417				
10884	2023-01-01 01:17:00	Peach Pit	Being so Normal	
2.205				
10885	2023-01-01 01:17:00	Kelly Clarkson	Santa, Can't You Hear Me	
0.278				

2.5 Finding top 10 favorite artists

Find the top **ten** artists with the highest total play time (in seconds). Plot your findings in a bar graph. (hint: start by creating a new DataFrame with only **artistName** and your time column. To proceed, you will also likely need the **groupby** command from Pandas.)

```
# Creating a new DataFrame with only artistName and secPlayed column
artist_playtime_df = df3[['artistName', 'secPlayed']]

# Grouping by artistName and sum play time
artist_total_playtime = artist_playtime_df.groupby('artistName').sum()


# Sorting artists based on total play time
sorted_artists = artist_total_playtime.sort_values(by='secPlayed',
ascending=False)

# Printing top ten artists
top_ten_artists = sorted_artists.head(10)

top_ten_artists
```

artistName	secPlayed
Ariana Grande	178996.003
Joji	137229.562
The Pretty Reckless	110293.430
Cage The Elephant	95587.575
Lana Del Rey	90543.113
Greta Van Fleet	88026.405
The Neighbourhood	85673.375
Gorillaz	84858.371
Led Zeppelin	77030.802
Arctic Monkeys	76444.236

```
# Plotting the Top 10 Songs
top_ten_artists.plot(kind='bar', figsize=(10, 6))
plt.title('Top 10 Artists by Total Play Time')
plt.xlabel('Artist')
plt.ylabel('Total Play Time (seconds)')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



2.6 Finding top 10 favorite songs

Find the top `ten` songs with the highest play time. Create a graph visualizing the results.

```
# Creating a new DataFrame with only trackName and secPlayed column
track_playtime_df = df3[['trackName', 'secPlayed']]

# Grouping by artistName and sum play time
track_total_playtime = track_playtime_df.groupby('trackName').sum()

# Sorting artists based on total play time
sorted_tracks = track_total_playtime.sort_values(by='secPlayed',
ascending=False)

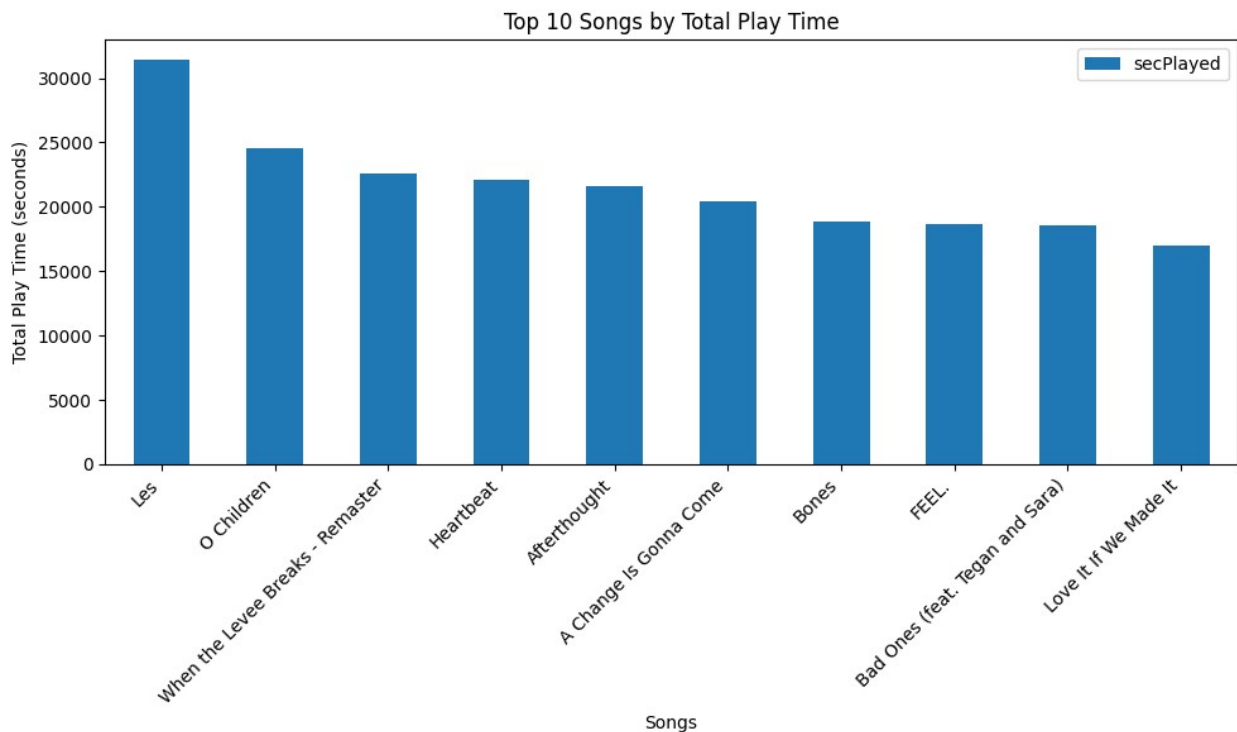
# Printing top ten artists
top_ten_tracks = sorted_tracks.head(10)

top_ten_tracks
```

	secPlayed
trackName	
Les	31403.364
0 Children	24558.414
When the Levee Breaks - Remaster	22631.721
Heartbeat	22056.629
Afterthought	21599.564
A Change Is Gonna Come	20414.317
Bones	18860.916


```
FEEL. 18696.637
Bad Ones (feat. Tegan and Sara) 18558.247
Love It If We Made It 17018.248
```

```
# Plotting the Top 10 Songs
top_ten_tracks.plot(kind='bar', figsize=(10, 6))
plt.title('Top 10 Songs by Total Play Time')
plt.xlabel('Songs')
plt.ylabel('Total Play Time (seconds)')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



Part 3: Further analysis

3.0 Average listening time by hour

Generate a plot that displays the average amount of time that music is played for each hour of the day.

```
# Creating a new DataFrame with only endTime and secPlayed column
music_play_df = df3[['endTime', 'secPlayed']]

# Extracting the hours from the endTime
music_play_df['endHour'] = music_play_df['endTime'].dt.hour
```

```

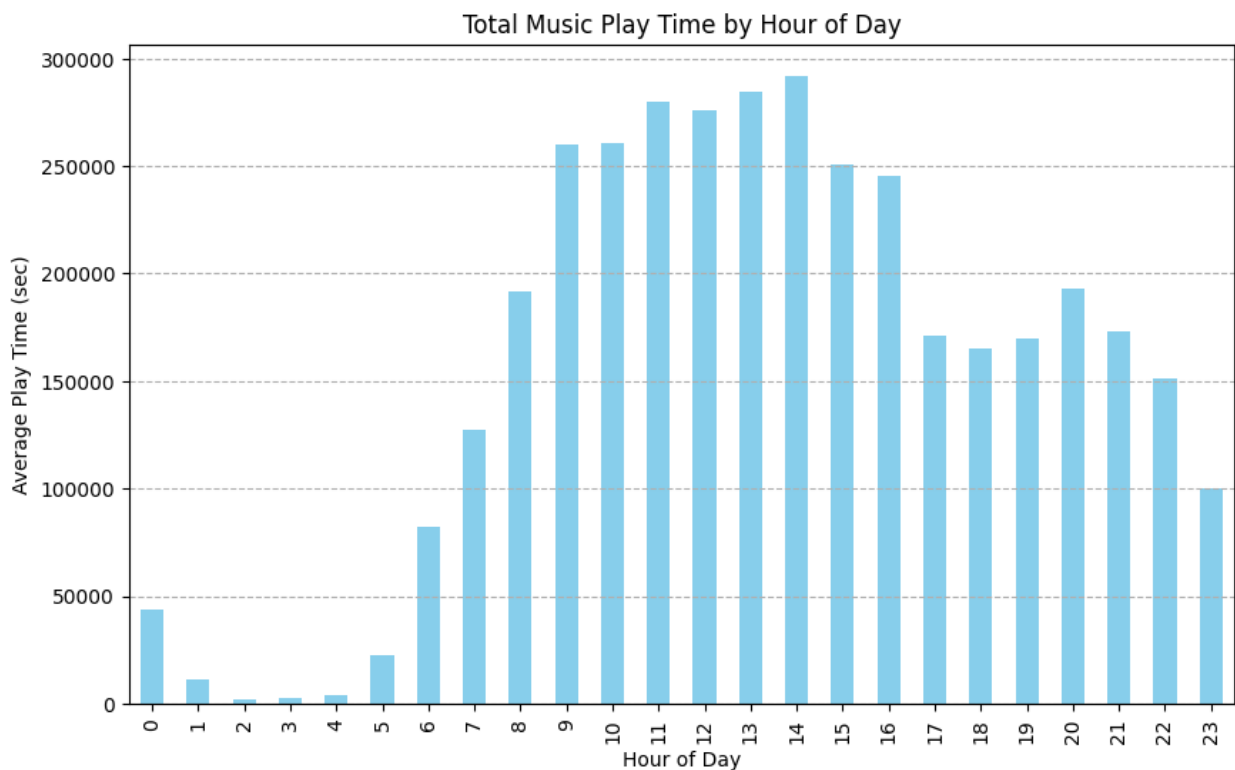
# Grouping by the hours and total the play time
total_play_time_per_hour = music_play_df.groupby('endHour')
['secPlayed'].sum()

# Grouping by the hours and average the play time
average_play_time_per_hour = music_play_df.groupby('endHour')
['secPlayed'].mean()

# Plot the result
plt.figure(figsize=(10, 6))
total_play_time_per_hour.plot(kind='bar', color='skyblue')
plt.title('Total Music Play Time by Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('Average Play Time (sec)')
plt.xticks(range(0, 24))
plt.grid(axis='y', linestyle='--')

# Show the plot
plt.show()

```




```

# Plot the result
plt.figure(figsize=(10, 6))
average_play_time_per_hour.plot(kind='bar', color='skyblue')
plt.title('Average Music Play Time by Hour of Day')
plt.xlabel('Hour of Day')

```

```
plt.ylabel('Average Play Time (sec)')
plt.xticks(range(0, 24))
plt.grid(axis='y', linestyle='--')

# Show the plot
plt.show()
```



Observations

Here is one interesting thing to notice if we use the the mean function while groupby it shows more average time in hour 0 to 6 although total listening hour is less on that time which can be seen from the previous graph and this is also normal that at that time the user of this data must be sleeping or resting. But the acverage come like this because as there are less skipped song on that time and the average here is shown by the total number of songs(both skipped and unskipped).


So, in the below I tried to deal with it by dividne by the total number days.

```
# Finding the total number days when the music is listened i.e. unique
days
days = music_play_df['endTime'].dt.day_of_year
unique_days = days.unique()
total_unique_days = len(unique_days)
```

```
# Getting average by dividing it by total number of days
mean_play_time_per_hour = total_play_time_per_hour/total_unique_days

# Plot the result
plt.figure(figsize=(10, 6))
mean_play_time_per_hour.plot(kind='bar', color='skyblue')
plt.title('Average Music Play Time by Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('Average Play Time (sec)')
plt.xticks(range(0, 24))
plt.grid(axis='y', linestyle='--')

# Show the plot
plt.show()
```



Now this look this look more understandable that how much average time the user has spent on listening in every hour.

3.1 Morning music and evening music

I think many people find that some types of music are more suitable for morning listening and some music is more suitable for evening listening. Create a plot that compares the play time of the artists *Leonard Cohen* and *Rage Against the Machine* on an hour-by-hour basis. See if there are any differences.

```
# Creating a new DataFrame with only Leonard Cohen and Rage Against The Machine
```

```
relevant_artists_df = df3[df3['artistName'].isin(['Leonard Cohen', 'Rage Against The Machine'])]
```

```
# Extracting the hours from the endTime
```

```
relevant_artists_df['hour'] = relevant_artists_df['endTime'].dt.hour
```

```
# Grouping by hour and artist, and sum the play time
```

```
hourly_playtime_comparison = relevant_artists_df.groupby(['hour', 'artistName'])['secPlayed'].sum().unstack()
```

```
/var/folders/2k/4yf6q3p91dn305kq0h_z85300000gn/T/
```

```
ipykernel_63153/2126971822.py:5: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy
```

```
relevant_artists_df['hour'] = relevant_artists_df['endTime'].dt.hour
```

```
hourly_playtime_comparison
```


```
artistName Leonard Cohen Rage Against The Machine
```

```
hour
```

0	486.593	2.024
1	19.065	NaN
4	271.905	17.242
5	371.150	NaN
6	190.543	28.610
7	296.875	12.516
8	595.872	42.402
9	829.501	226.606
10	2068.587	335.398
11	1657.468	813.869
12	1078.065	334.003
13	1425.191	238.100
14	1731.914	141.379
15	1418.022	276.347
16	1142.896	1577.069
17	475.653	1183.734
18	1545.162	913.117
19	2101.166	332.719
20	2978.077	123.206
21	2028.490	208.668
22	619.896	234.753
23	121.525	104.835

```
# Plotting the results
plt.figure(figsize=(10, 6))
hourly_playtime_comparison.plot(kind='line', marker='o',
linestyle='-')
plt.title('Hourly Play Time Comparison: Leonard Cohen vs. Rage Against
The Machine')
plt.xlabel('Hour of the Day')
plt.ylabel('Total Play Time (seconds)')
plt.xticks(range(24))
plt.grid(True)
plt.legend(title='Artist')
plt.tight_layout()
plt.show()
```

<Figure size 1000x600 with 0 Axes>



3.2 Analysing skipped songs

Determining whether a song was skipped or listened to can be challenging. For this analysis, we'll simplify by defining a skipped song as any track played for less than 30 seconds. Conversely, a song played for 30 seconds or more is considered listened to. Add a column to your DataFrame to reflect this criteria: set the value to 1 if the song was played for less than 30 seconds (indicating a skipped song), and 0 if it was played for 30 seconds or longer.

```
# Adding new column skipped songs

skippedTrack = []

for i in df3['secPlayed']:
    if i < 30 :
        skippedTrack.append(1)
    else:
        skippedTrack.append(0)

df3.insert(4, 'skippedTrack', skippedTrack)

df3.head()
```

		endTime	artistName	trackName \
10881	2023-01-01	01:17:00	Ariana Grande	7 rings
10882	2023-01-01	01:17:00	Ariana Grande	7 rings
10883	2023-01-01	01:17:00	Ariana Grande	positions
10884	2023-01-01	01:17:00	Peach Pit	Being so Normal
10885	2023-01-01	01:17:00	Kelly Clarkson	Santa, Can't You Hear Me

	secPlayed	skippedTrack
10881	0.139	1
10882	0.487	1
10883	0.417	1
10884	2.205	1
10885	0.278	1


3.3 Plotting skipped songs

Create a pie-chart that compares amount of skipped songs to amount of non-skipped songs.

```
# Step 1: Group the DataFrame by the 'skipped' column and count the
occurrences
skipped_counts = df3.groupby('skippedTrack').size()

# Step 2: Plot the results in a pie chart
plt.figure(figsize=(8, 6))
plt.pie(skipped_counts, labels=['Non-Skipped', 'Skipped'],
autopct='%1.1f%%', startangle=140, colors=['lightblue', 'lightcoral'])
plt.title('Comparison of Skipped vs. Non-Skipped Songs')
plt.axis('equal') # Equal aspect ratio ensures that pie is drawn as a
circle.
plt.show()
```

Comparison of Skipped vs. Non-Skipped Songs



3.4 Artists by percentage of songs skipped

For each artist in the dataset, calculate which percentage of their songs was skipped. Store this information in a new DataFrame called `df_skipped`. Store the percentage of skipped songs in a new column named `SkipRate` **Example:** If an artist has **100** songs in your dataset and **25** of these were skipped, the percentage of skipped songs for this artist would be $\frac{25}{100} = 25\%$

```
# Grouping the DataFrame by the 'artistName' column
grouped_by_artist = df3.groupby('artistName')

# Calculating the total number of songs and the number of skipped
songs for each artist
total_songs = grouped_by_artist.size()
skipped_songs = grouped_by_artist['skippedTrack'].sum()

# computing the percentage of skipped songs for each artist
skip_rate = (skipped_songs / total_songs) * 100

# Creating a new DataFrame to store the percentage of skipped songs
```



```

for each artist
df_skipped = pd.DataFrame({'TotalSongs': total_songs, 'SkippedTrack':
skipped_songs, 'SkipRate': skip_rate})

# Display the new DataFrame
df_skipped

```

artistName	TotalSongs	SkippedTrack	SkipRate
10cc	28	19	67.857143
2Pac	513	442	86.159844
3 Doors Down	2	1	50.000000
4 Non Blondes	122	88	72.131148
50 Cent	28	19	67.857143
...
squeeda	3	2	66.666667
tenkousei.	37	37	100.000000
trxxshed	2	1	50.000000
xander.	8	3	37.500000
Édith Piaf	155	146	94.193548

[956 rows x 3 columns]

3.5 Comparing artists by skip-rate

Find the **three** top artists with the lowest skip-rate and the **three** with the highest. Print their names, along with their skip-rate.

```

# Finding three top artists with lowest skip-rate
print(df_skipped['SkipRate'].nsmallest(3))

artistName
Gloria Gaynor      0.000000
Roc Boyz          11.111111
LACES             14.285714
Name: SkipRate, dtype: float64

# Finding three top artists with highest skip-rate
print(df_skipped['SkipRate'].nlargest(3))

artistName
A Problem Squared    100.0
Acid Ghost           100.0
Albert Hammond Jr    100.0
Name: SkipRate, dtype: float64

```

Part 4: God Is a Data Scientist - The Ariana Deep-Dive

4.0 Ariana-DataFrame:

Create a new DataFrame called `df_ariana`, containing only rows with music by Ariana Grande.

```
# Creating a dataframe with only rows with music by Ariana Grande
df_ariana = df3[df3['artistName'] == 'Ariana Grande']
```

df_ariana

	trackName \	endTime	artistName	
10881	2023-01-01 01:17:00	Ariana Grande	7	
10882	2023-01-01 01:17:00	Ariana Grande	7	
10883	2023-01-01 01:17:00	Ariana Grande		
10887	2023-01-01 01:17:00	Ariana Grande	Santa	
10888	2023-01-01 01:17:00	Ariana Grande	Right There (feat. Big Sean)	
...	
167415	2023-12-07 17:46:00	Ariana Grande	Almost Is Never	
167422	2023-12-07 20:51:00	Ariana Grande		
167428	2023-12-07 21:13:00	Ariana Grande	pete	
167435	2023-12-07 21:13:00	Ariana Grande	off the table (with The Weeknd)	
167436	2023-12-07 21:14:00	Ariana Grande	my	

	secPlayed	skippedTrack
10881	0.139	1
10882	0.487	1
10883	0.417	1
10887	12.293	1
10888	22.929	1
...
167415	28.483	1
167422	26.220	1
167428	0.603	1
167435	13.448	1
167436	23.757	1

[19337 rows x 5 columns]

```
df_ariana.info()

<class 'pandas.core.frame.DataFrame'>
Index: 19337 entries, 10881 to 167436
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   endTime         19337 non-null  datetime64[ns]
1   artistName      19337 non-null  object
2   trackName       19337 non-null  object
3   secPlayed       19337 non-null  float64
4   skippedTrack    19337 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(1), object(2)
memory usage: 906.4+ KB

df3.info()

<class 'pandas.core.frame.DataFrame'>
Index: 156539 entries, 10881 to 167438
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   endTime         156539 non-null  datetime64[ns]
1   artistName      156539 non-null  object
2   trackName       156539 non-null  object
3   secPlayed       156539 non-null  float64
4   skippedTrack    156539 non-null  int64
dtypes: datetime64[ns](1), float64(1), int64(1), object(2)
memory usage: 7.2+ MB
```


4.1 Average skip rate

Create a histogram of the distribution of the skip-rate values of the different artists in your DataFrame `df_skipped`, with skip rates on one axis and number of artists on the other. Then, retrieve the skip rate for Ariana Grande from your DataFrame `df_skipped`. Run the code in the cell below. Where on this distribution does Ariana Grande fall? Do I skip her songs more than average, or less?

```
# Plotting a histogram of skip-rate values
plt.figure(figsize=(10, 6))
plt.hist(df_skipped['SkipRate'], bins=20, color='skyblue',
         edgecolor='black')
plt.title('Distribution of Skip-Rate Values')
plt.xlabel('Skip Rate (%)')
plt.ylabel('Number of Artists')
plt.grid(True)
plt.show()

# Retrieving the skip rate for Ariana Grande
```

```
ariana_grande_skip_rate = df_skipped.loc['Ariana Grande', 'SkipRate']
print("Skip rate for Ariana Grande:", ariana_grande_skip_rate)
```



```
Skip rate for Ariana Grande: 99.52939959662822
```

```
count_skipped_artist_df = df_skipped[df_skipped['SkipRate'] >= 90]
```

```
count_skipped_artist_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 181 entries, A Problem Squared to Édith Piaf
```

```
Data columns (total 3 columns):
```

#	Column	Non-Null Count	Dtype
0	TotalSongs	181 non-null	int64
1	SkippedTrack	181 non-null	int64
2	SkipRate	181 non-null	float64

```
dtypes: float64(1), int64(2)
```

```
memory usage: 5.7+ KB
```

```
df_skipped.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 956 entries, 10cc to Édith Piaf
```

```
Data columns (total 3 columns):
```

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

```
---  -----  -----  ---
0    TotalSongs    956 non-null    int64
1    SkippedTrack  956 non-null    int64
2    SkipRate      956 non-null    float64
dtypes: float64(1), int64(2)
memory usage: 62.2+ KB
```

Part 4: Questions

Q1: Did I skip a lot of Ariana Grande's songs, or did I not, compared to the rest of the dataset?

Answer1: As we can see skip rate is close to 100%, so a lot of songs of Ariana Grande is Skipped.

Compare to the rest of the dataset we can see that 20% (181 out 956) of the artist's song

skipped more than 80%. Q2: What might be some possible reasons for Ariana Grande to be my nr.1 artist? Answer2: In the dataset of 2023 spotify data we can see out of 156539 tracks there are 19337 tracks of Ariana Grande. And so many of songs are skipped so that makes Ariana the no. 1 artist.