

# CA3 - Lumifruit Edibility Classification

Jony Karmakar

## Imports

```
# Importing Libraries
# =====
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.impute import SimpleImputer

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```

## Reading data

```
# Loading the dataset
# =====
df_train = pd.read_csv('train.csv', index_col=0)
df_test = pd.read_csv('test.csv', index_col=0)
```

## Data exploration and visualisation

```
#Checking for missing data
# =====
null_counts = df_train.isnull().sum().sum()
print(f'Total Missing values: {null_counts}' )
```

Total Missing values: 37

```
# Checking the info
df_train.info()
```

<class 'pandas.core.frame.DataFrame'>

Index: 1248 entries, 0 to 1247

Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	Acoustic Firmness Index	1245 non-null	float64
1	Atmospheric Pressure at Harvest (Pa)	1248 non-null	float64
2	Bitterness Scale	1244 non-null	float64

3	Circumference (mm)	1246	non-null	float64
4	Color Intensity (a.u.)	1248	non-null	float64
5	Find Distance from Main Vulcano (km)	1246	non-null	float64
6	Length (mm)	1244	non-null	float64
7	Luminescence Intensity (a.u.)	1247	non-null	float64
8	Magnetic orientation (degree)	1245	non-null	float64
9	Odor index (a.u.)	1246	non-null	float64
10	Seed Count	1247	non-null	float64
11	Skin Thickness (mm)	1247	non-null	float64
12	Soil pH where Grown	1245	non-null	float64
13	Sugar Content (mg)	1245	non-null	float64
14	Weight (mg)	1244	non-null	float64
15	pH	1244	non-null	float64
16	Edible	1248	non-null	float64
dtypes: float64(17)				
memory usage: 175.5 KB				
df_train.describe()				
Acoustic Firmness Index    Atmospheric Pressure at Harvest				
(Pa) \				
count	1245.000000		1248.000000	
mean	21.570077		101327.543269	
std	8.131888		4772.582203	
min	0.600000		83825.000000	
25%	17.300000		98095.750000	
50%	21.300000		101357.000000	
75%	25.300000		104470.750000	
max	156.520701		115636.000000	
Bitterness Scale    Circumference (mm)    Color Intensity (a.u.) \				
count	1244.000000	1246.000000	1248.000000	
mean	1.808682	211.046062	63.300962	
std	1.960279	24.652278	7.118135	
min	0.000000	145.867667	46.060000	
25%	0.000000	192.982222	58.577500	
50%	1.000000	212.429730	62.390000	
75%	3.000000	226.313333	67.202500	
max	7.000000	274.719407	97.810000	
Find Distance from Main Vulcano (km)    Length (mm) \				
count	1246.000000	1244.000000		
mean	44.188204	33.589646		

std	28.331532	3.874355
min	2.550790	23.451799
25%	21.203453	30.777936
50%	38.880858	33.784401
75%	60.680816	36.049083
max	109.952511	43.691515

	Luminescence Intensity (a.u.)	Magnetic orientation (degree) \
count	1247.000000	1245.000000
mean	0.002850	178.933737
std	0.003607	102.113321
min	0.000004	0.085357
25%	0.000627	90.582930
50%	0.001757	180.043851
75%	0.003767	266.481765
max	0.044624	359.443812

	Odor index (a.u.)	Seed Count	Skin Thickness (mm) \
count	1246.000000	1247.000000	1247.000000
mean	16.602442	48.926812	3.839615
std	14.882420	35.871733	1.483029
min	1.722825	0.000000	1.000000
25%	7.615908	15.852008	3.000000
50%	12.258785	47.932723	4.000000
75%	20.302048	80.413165	5.000000
max	152.041780	112.968004	8.500000

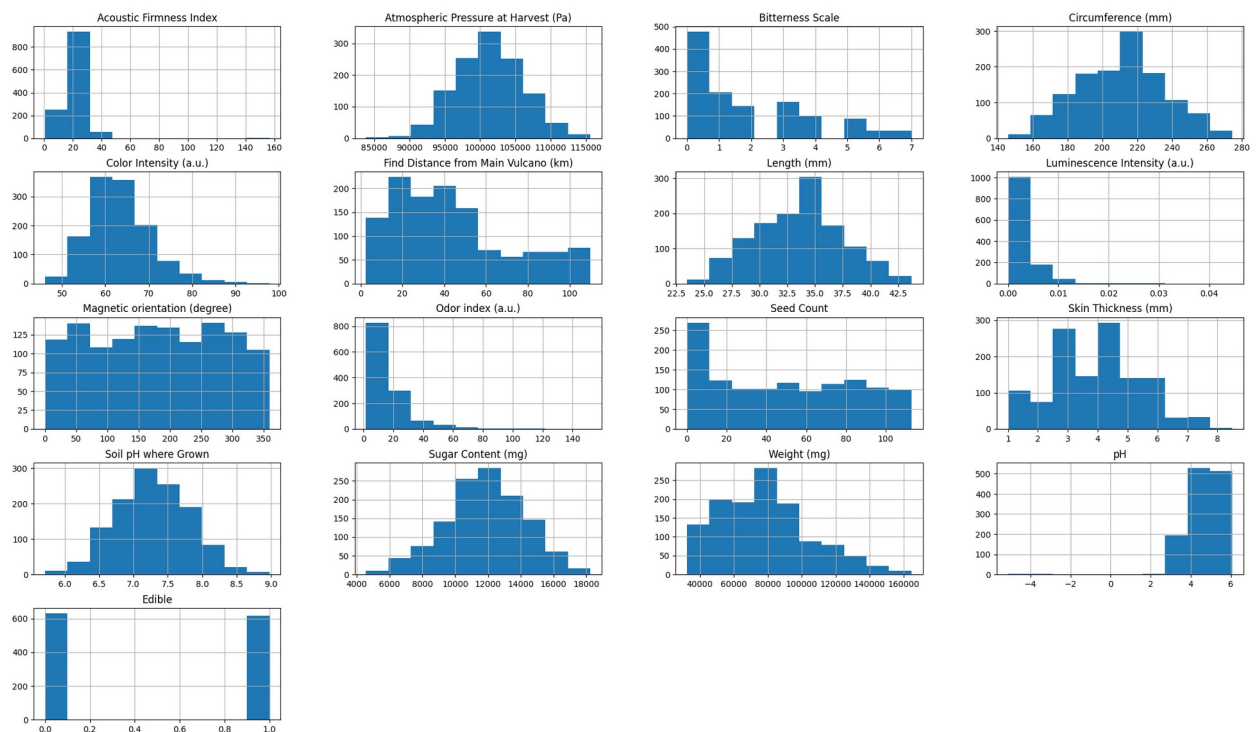
	Soil pH where Grown	Sugar Content (mg)	Weight (mg)
pH \			
count	1245.000000	1245.000000	1244.000000
1244.000000			
mean	7.278739	11835.978313	78446.779224
4.614780			
std	0.529815	2417.733775	26637.766132
0.993407			
min	5.710000	4566.000000	32352.182600
5.120000			
25%	6.910000	10353.000000	57087.682600
4.174979			
50%	7.280000	11807.000000	77440.682600
4.910000			
75%	7.660000	13534.000000	92664.102600
5.080000			
max	8.980000	18246.000000	164679.342600
6.081918			

	Edible
count	1248.000000
mean	0.494391
std	0.500169

min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

## Histogram

```
# Plotting histograms for all columns in the dataset
df_train.hist(bins=10, figsize=(26, 15))
plt.show()
```



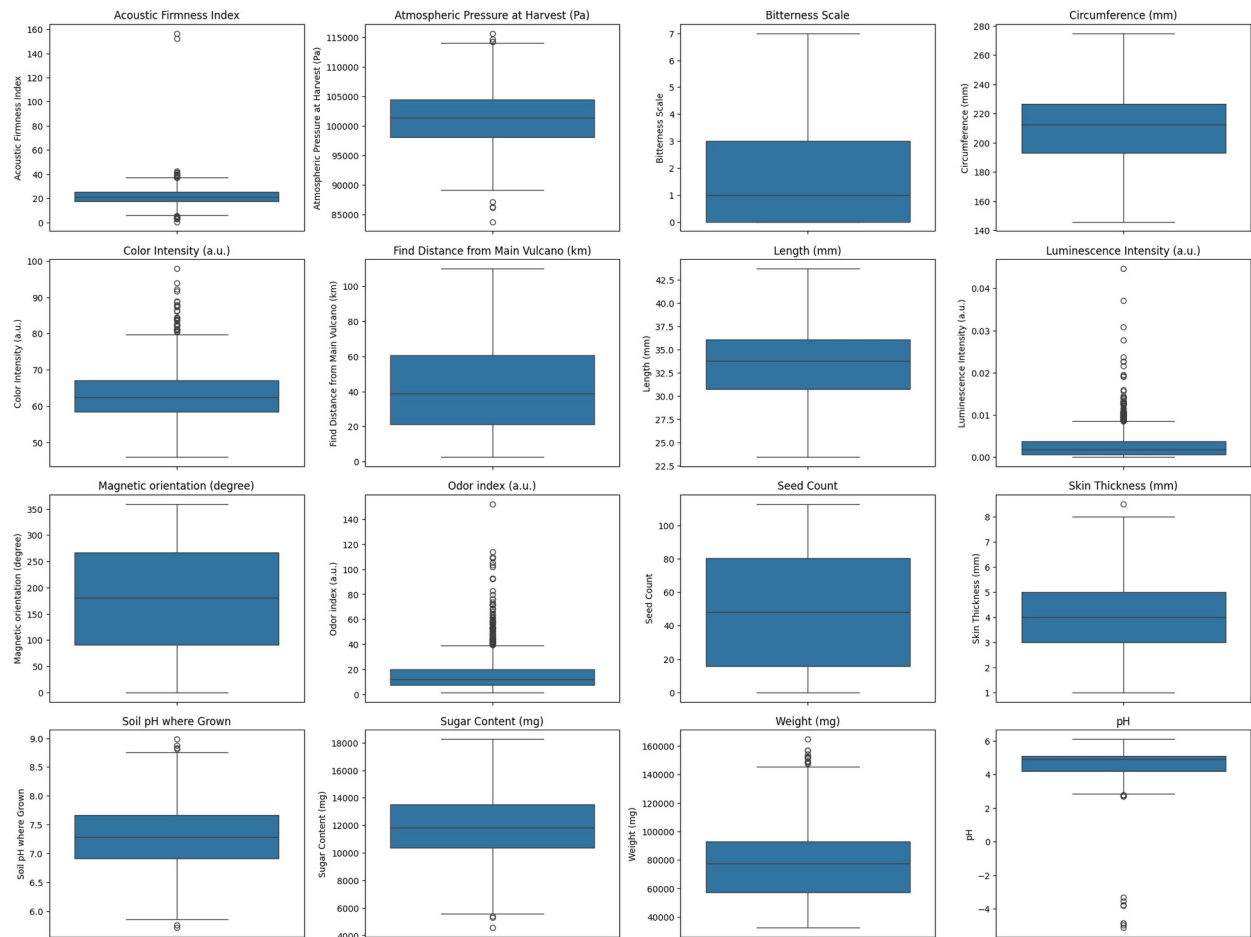
## Box Plots

```
# Removing the target variable
features = df_train.columns[:-1]

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.boxplot(y=df_train[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



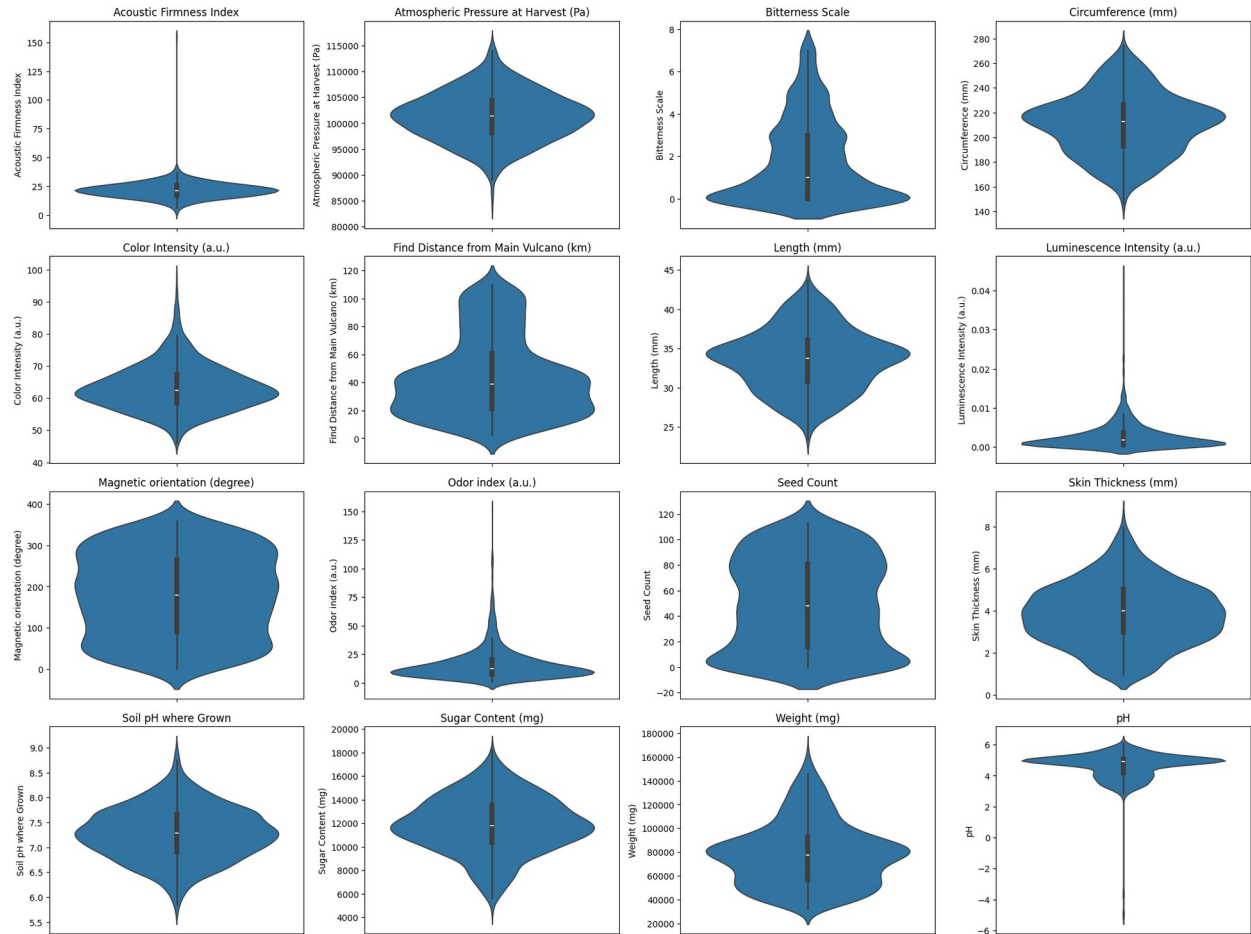
## Violin Plots

```
# Removing the target variable
features = df_train.columns[:-1]

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.violinplot(y=df_train[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```




## Heatmap

```
corr_matrix = df_train.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, cbar=True)

plt.show()
```



## Scatter Plot

```
sns.pairplot(df_train, hue="Edible")
plt.show()
```







```

df_train_imp = pd.DataFrame(imputed_data, columns=df_train.columns)
df_train_imp.index = df_train.index

df_train_imp.isnull().sum()

```

Acoustic Firmness Index	0
Atmospheric Pressure at Harvest (Pa)	0
Bitterness Scale	0
Circumference (mm)	0
Color Intensity (a.u.)	0
Find Distance from Main Vulcano (km)	0
Length (mm)	0
Luminescence Intensity (a.u.)	0
Magnetic orientation (degree)	0
Odor index (a.u.)	0
Seed Count	0
Skin Thickness (mm)	0
Soil pH where Grown	0
Sugar Content (mg)	0
Weight (mg)	0
pH	0
Edible	0
dtype: int64	

## Detecting the Outliers

```

# Detect the outliers using the z-score
outliers_count = {}

# Iterating through each column in df_train_dna
for column in df_train_imp.columns:
    samples = df_train_imp[column].values
    z_scores = (samples - np.mean(samples)) / np.std(samples)
    outliers = np.abs(z_scores) > 3
    outliers_count[column] = np.sum(outliers)

for feature, count in outliers_count.items():
    print(f"Number of outliers in {feature}: {count}")

```

```

Number of outliers in Acoustic Firmness Index: 2
Number of outliers in Atmospheric Pressure at Harvest (Pa): 3
Number of outliers in Bitterness Scale: 0
Number of outliers in Circumference (mm): 0
Number of outliers in Color Intensity (a.u.): 12
Number of outliers in Find Distance from Main Vulcano (km): 0
Number of outliers in Length (mm): 0
Number of outliers in Luminescence Intensity (a.u.): 18
Number of outliers in Magnetic orientation (degree): 0
Number of outliers in Odor index (a.u.): 28
Number of outliers in Seed Count: 0

```

```
Number of outliers in Skin Thickness (mm): 1
Number of outliers in Soil pH where Grown: 2
Number of outliers in Sugar Content (mg): 1
Number of outliers in Weight (mg): 1
Number of outliers in pH: 8
Number of outliers in Edible: 0
```

For this assignment outliers have been ignored.

## Data preprocessing and visualisation

### Splitting the Dataset

```
#-----
# Splitting the dataset into test and train
#-----
X_imp = df_train_imp.iloc[:, :-1].values
y_imp = df_train_imp.iloc[:, -1].values
X_train_imp, X_test_imp, y_train_imp, y_test_imp =
train_test_split(X_imp, y_imp, test_size=0.20, stratify=y_imp,
random_state=42)
```

### Scaling the Dataset

```
#
=====
# Standardising using the StandardScaler (Taken from the lecture also
this is the common way of scaling)
#
=====
# Initialise standard scaler and compute mean and stddev from training
data
sc = StandardScaler()
sc.fit(X_train_imp)

# Transform (standardise) both X_train_imp and X_test_imp with mean
and stddev from
# training data
X_train_imp_sc = sc.transform(X_train_imp)
X_test_imp_sc = sc.transform(X_test_imp)
```

### Exploring training data after cleaning and scaling

```
# Converting numpy array to dataframe
X_train_imp_sc_df = pd.DataFrame(X_train_imp_sc)
X_test_imp_sc_df = pd.DataFrame(X_test_imp_sc)

X_train_imp_sc_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 998 entries, 0 to 997
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	0	998 non-null	float64
1	1	998 non-null	float64
2	2	998 non-null	float64
3	3	998 non-null	float64
4	4	998 non-null	float64
5	5	998 non-null	float64
6	6	998 non-null	float64
7	7	998 non-null	float64
8	8	998 non-null	float64
9	9	998 non-null	float64
10	10	998 non-null	float64
11	11	998 non-null	float64
12	12	998 non-null	float64
13	13	998 non-null	float64
14	14	998 non-null	float64
15	15	998 non-null	float64

```
dtypes: float64(16)
```

```
memory usage: 124.9 KB
```

```
X_train_imp_sc_df.describe()
```

	0	1	2	3
4 \				
count	9.980000e+02	9.980000e+02	9.980000e+02	9.980000e+02
	9.980000e+02			
mean	1.203224e-15	7.591345e-16	-1.690921e-17	-5.385138e-15
	7.671441e-16			
std	1.000501e+00	1.000501e+00	1.000501e+00	1.000501e+00
	1.000501e+00			
min	-2.790267e+00	-3.192178e+00	-9.400674e-01	-2.637676e+00
	2.421179e+00			
25%	-5.767442e-01	-6.614693e-01	-9.400674e-01	-7.379985e-01
	6.535684e-01			
50%	-1.002903e-02	5.589991e-03	-4.263936e-01	5.209417e-02
	1.180030e-01			
75%	5.033483e-01	6.355494e-01	6.009541e-01	6.287116e-01
	5.392977e-01			
max	1.800094e+01	2.812833e+00	2.655650e+00	2.558980e+00
	4.819973e+00			
5		6	7	8
9 \				
count	9.980000e+02	9.980000e+02	9.980000e+02	9.980000e+02
	9.980000e+02			
mean	-3.150453e-16	-2.669875e-16	-9.433558e-17	2.036225e-15

```

1.459532e-15
std      1.000501e+00  1.000501e+00  1.000501e+00  1.000501e+00
1.000501e+00
min      -1.485719e+00 -2.620729e+00 -8.202004e-01 -1.753441e+00 -
9.915981e-01
25%      -8.184595e-01 -7.314885e-01 -6.379509e-01 -8.690324e-01 -
6.032569e-01
50%      -1.719907e-01  4.262328e-02 -3.115349e-01  1.105358e-02 -
3.013487e-01
75%       6.343581e-01  6.444112e-01  2.732684e-01  8.434427e-01
2.636343e-01
max       2.290280e+00  2.591265e+00  1.208737e+01  1.793447e+00
8.919630e+00

```

```

                                10          11          12          13
14 \
count  9.980000e+02  9.980000e+02  9.980000e+02  9.980000e+02
9.980000e+02
mean   9.682747e-16  1.779917e-17 -1.940109e-15 -7.600244e-16 -
3.287150e-14
std     1.000501e+00  1.000501e+00  1.000501e+00  1.000501e+00
1.000501e+00
min    -1.400326e+00 -1.877736e+00 -2.908752e+00 -3.057209e+00 -
1.686056e+00
25%    -9.239207e-01 -8.737602e-01 -7.038083e-01 -6.294217e-01 -
8.081063e-01
50%    -3.176661e-02  1.302153e-01  6.258341e-03 -5.000689e-03 -
4.853229e-02
75%     8.728676e-01  7.995323e-01  7.350109e-01  6.914323e-01
5.476844e-01
max     1.766415e+00  3.142142e+00  2.996013e+00  2.497759e+00
3.197890e+00

```

```

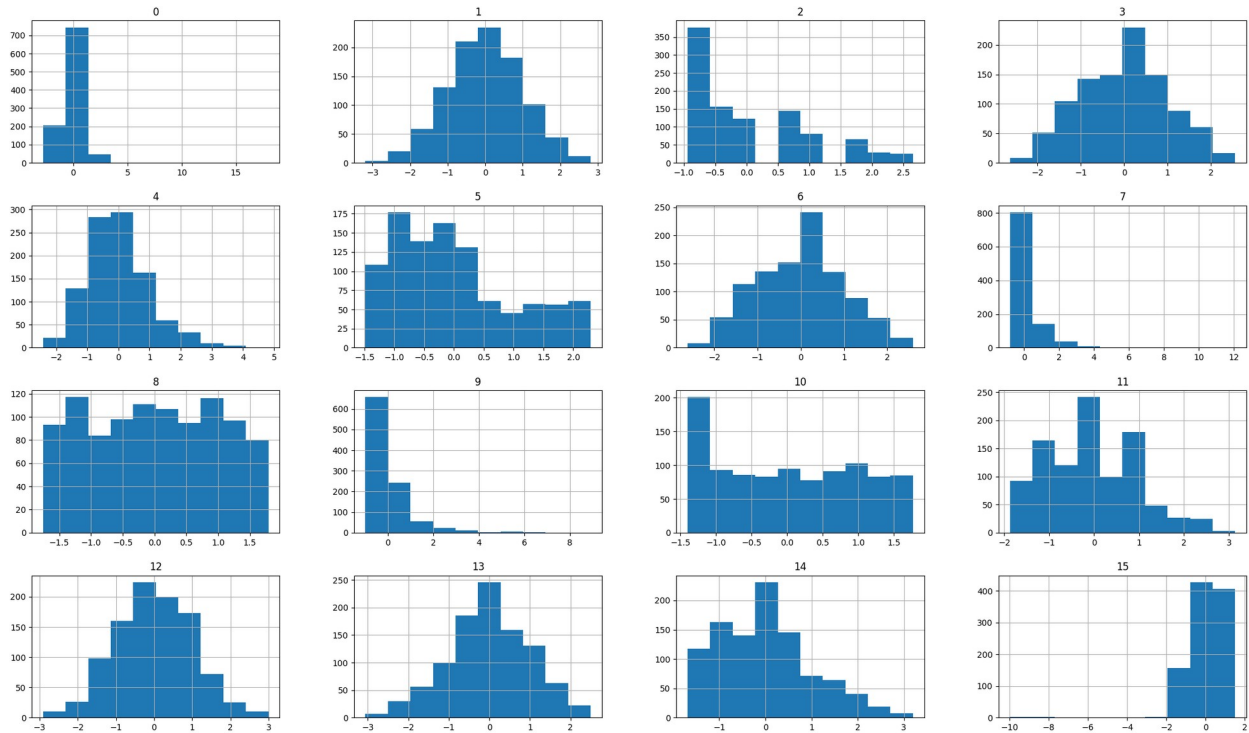
                                15
count  9.980000e+02
mean   7.977587e-15
std     1.000501e+00
min    -9.988594e+00
25%    -4.590113e-01
50%     2.962997e-01
75%     4.707939e-01
max     1.509464e+00

```

```

# Plotting histograms for all columns in the dataset
X_train_imp_sc_df.hist(bins=10, figsize=(26, 15))
plt.show()

```



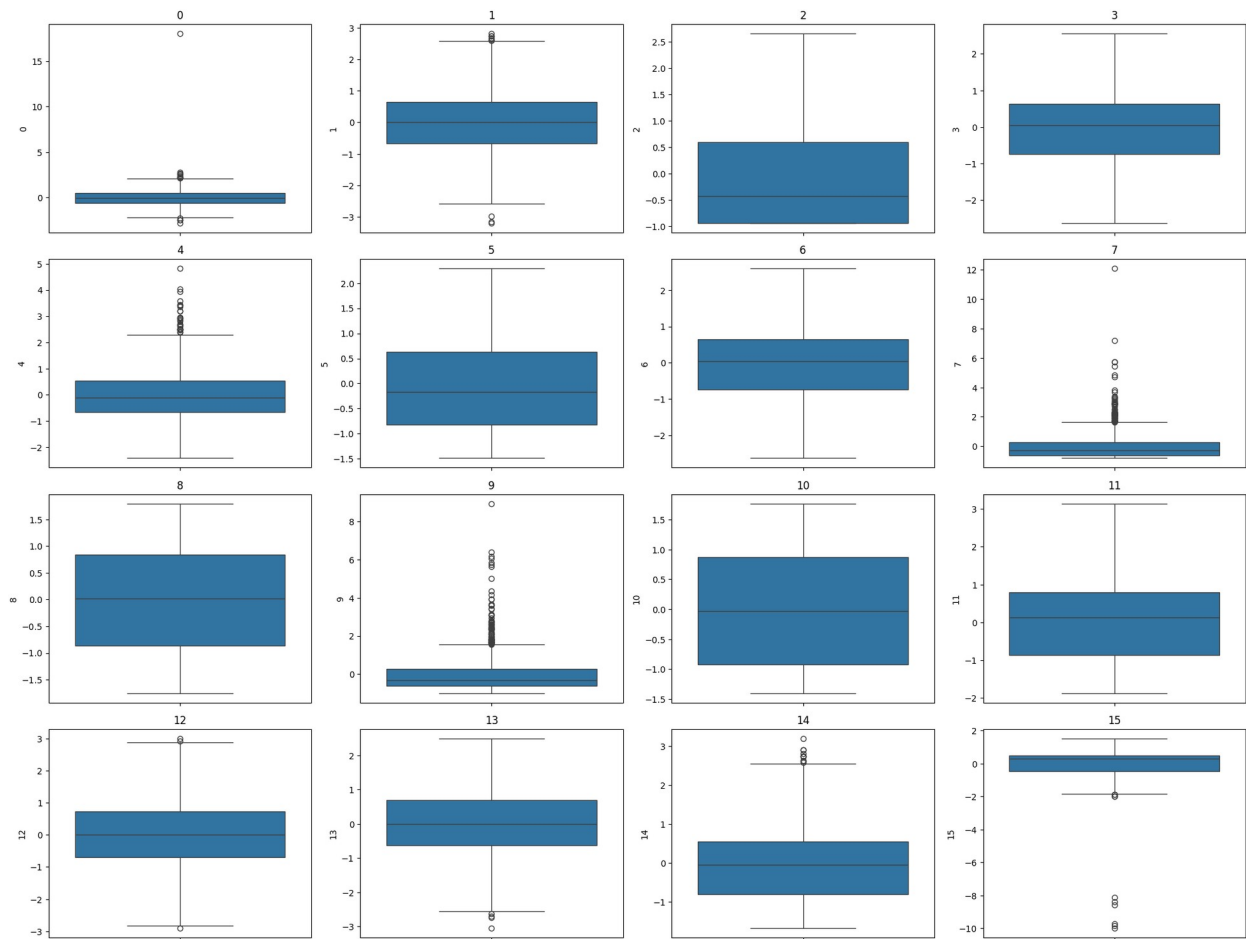
## Box Plots

```
features = X_train_imp_sc_df.columns

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.boxplot(y=X_train_imp_sc_df[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



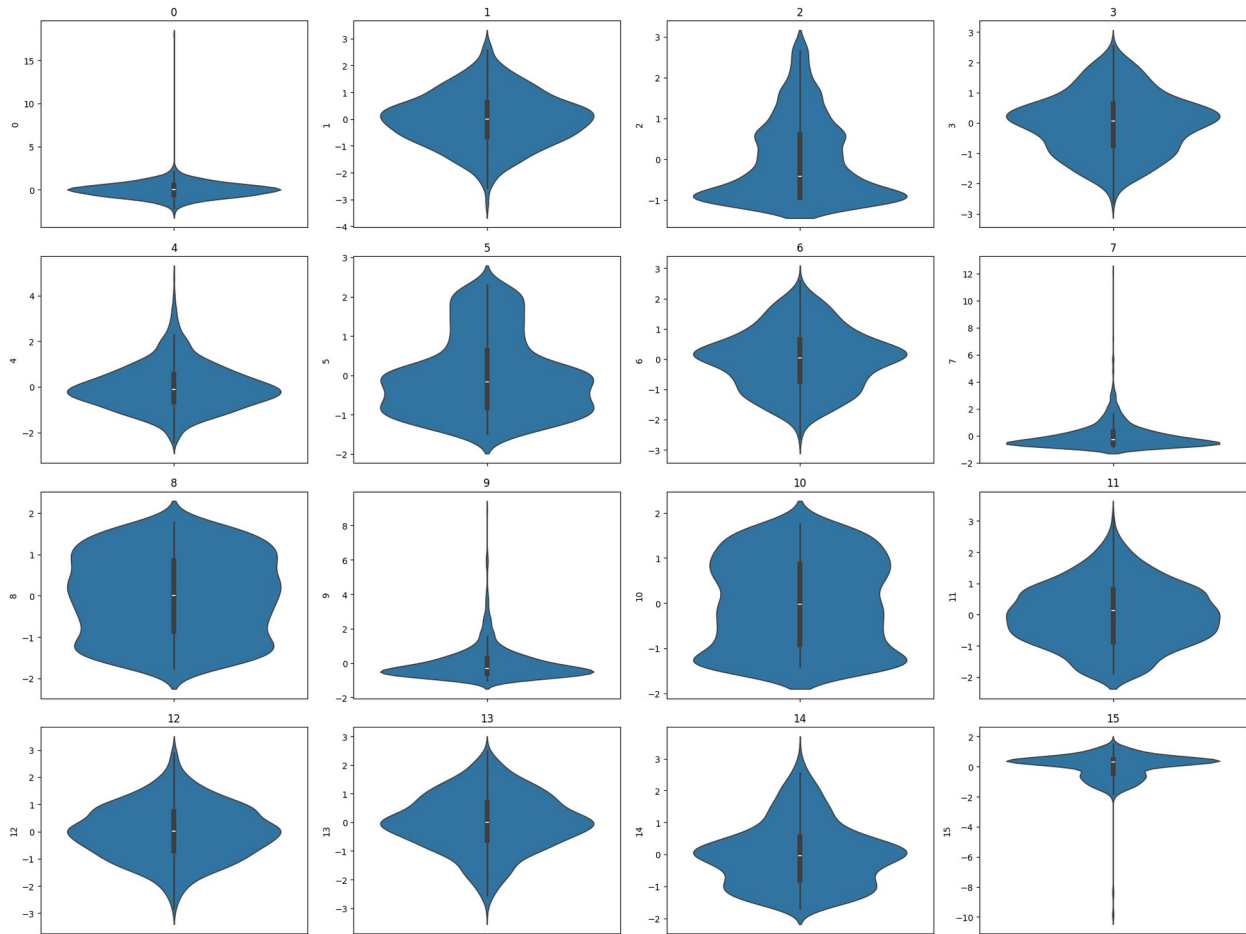
## Violin Plots

```
features = X_train_imp_sc_df.columns

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.violinplot(y=X_train_imp_sc_df[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```



## Modelling

### Random Forest

```
#-----
#-----
# RANDOM FOREST()
# Parameters tuned for this model
# 'n_estimators' : Number of Trees(High->Robustness but time
consuming)
# 'max_depth' : Max depth of Trees(High->Accurate model but
overfitting)
# 'max_features' : The number of features to consider when looking for
the best split
#-----
#-----

# Parameters to test
param_grid = {'n_estimators': [50, 100, 200], 'max_depth': [None, 5,
10, 25], 'max_features': [*np.arange(0.1, 1.1, 0.1)]}

# Store the best parameter combination and its accuracy
```



```

best_params_rf = None
best_accuracy_rf = 0

for n_estimators in param_grid['n_estimators']:
    for max_depth in param_grid['max_depth']:
        for max_feature in param_grid['max_features']:
            accuracies = []
            for random_state in range(40, 50):
                # Splitting Dataset
                X_train, X_val, y_train, y_val =
train_test_split(X_train_imp, y_train_imp, test_size=0.25,
random_state=random_state)
                # Standardize the data
                scaler = StandardScaler()
                X_train_scaled = scaler.fit_transform(X_train)
                X_val_scaled = scaler.transform(X_val)

                # Defining the model
                rfc =
RandomForestClassifier(n_estimators=n_estimators, max_depth=max_depth,
max_features=max_feature, random_state=42)
                # Fitting the model
                rfc.fit(X_train_scaled, y_train)
                # Evaluating on the test dataset
                y_pred = rfc.predict(X_val_scaled)
                accuracy = accuracy_score(y_val, y_pred)
                accuracies.append(accuracy)

            # Measuring the average accuracy for different splits of
dataset
            accuracy = np.mean(accuracies)

            if accuracy > best_accuracy_rf:
                best_accuracy_rf = accuracy
                best_params_rf= {'n_estimators': n_estimators,
'max_depth': max_depth, 'max_features': max_feature}

print("Best Parameters:", best_params_rf)
print("Best Accuracy:", best_accuracy_rf)

Best Parameters: {'n_estimators': 100, 'max_depth': 10,
'max_features': 1.0}
Best Accuracy: 0.8976000000000001

```

Testing for X\_train and X\_test of Train\_df

```

rfc =
RandomForestClassifier(n_estimators=best_params_rf['n_estimators'],
max_depth=best_params_rf['max_depth'],
max_features=best_params_rf['max_features'], random_state=42)

```

```

rfc.fit(X_train_imp_sc, y_train_imp)

y_pred_rfc = rfc.predict(X_test_imp_sc)
accuracy_rfc = accuracy_score(y_pred_rfc, y_test_imp)
print(f'For Random Forest with Best Parameters: {best_params_rf} and
Accuracy: {accuracy_rfc}')

For Random Forest with Best Parameters: {'n_estimators': 100,
'max_depth': 10, 'max_features': 1.0} and Accuracy: 0.904

```

Random Forest

For Random Forest we can see it has the best accuracy

## Final evaluation

```

#-----
# Train the best model on the entire training set
#-----
scaler = StandardScaler()
scaler.fit(X_imp)
X_train_scaled = scaler.transform(X_imp)
rfc.fit(X_train_scaled, y_imp)

#-----
# Use the test set to evaluate the best k on unseen data
#-----
X_test_scaled = scaler.transform(df_test)
y_pred_rf = rfc.predict(X_test_scaled)

```

## Kaggle submission

```

#-----
# Creating csv file for Kaggle submission
#-----
DF = pd.DataFrame(y_pred_rf)
headers = ["Edible"]
DF.columns = headers
DF.to_csv('predictions.csv', index_label='index', index=True)

```