

CA4 - Liver Disease Prediction

Jony Karmakar

Imports

```
# Importing Libraries
# =====
import matplotlib.pyplot as plt
from sklearn.metrics import ConfusionMatrixDisplay
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.impute import SimpleImputer

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

from sklearn.decomposition import PCA
from sklearn.model_selection import GridSearchCV
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.metrics import confusion_matrix
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
as LDA

from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_predict
```

Reading data

```
# Loading the dataset
# =====
df_train = pd.read_csv('train.csv', index_col=0)
df_test = pd.read_csv('test.csv', index_col=0)
```

Data exploration and visualisation

```
#Checking for missing data
# =====
```

```
null_counts = df_train.isnull().sum().sum()
print(f'Total Missing values: {null_counts}' )
```

Total Missing values: 0

```
# Checking the info of the dataset
```

```
# =====
```

```
df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 703 entries, 0 to 702
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	index	703 non-null	int64
1	AFP (ng/mL)	703 non-null	float64
2	ALP (U/L)	703 non-null	float64
3	ALT (U/L)	703 non-null	float64
4	AST (U/L)	703 non-null	float64
5	Age	703 non-null	int64
6	Albumin (g/dL)	703 non-null	float64
7	Alcohol_Use (yes/no)	703 non-null	object
8	Bilirubin (mg/dL)	703 non-null	float64
9	CRP (mg/L)	703 non-null	float64
10	Diabetes (yes/no)	703 non-null	object
11	Fibroscan (kPa)	703 non-null	float64
12	GGT (U/L)	703 non-null	float64
13	Gender	703 non-null	object
14	Hemoglobin (g/dL)	703 non-null	float64
15	IL-6 (pg/mL)	703 non-null	float64
16	Obesity (yes/no)	703 non-null	object
17	PT/INR	703 non-null	float64
18	Platelets (10 ⁹ /L)	703 non-null	float64
19	RBC (10 ¹² /L)	703 non-null	float64
20	Serum_Ammonia (μmol/L)	703 non-null	float64
21	Serum_Copper (μg/dL)	703 non-null	float64
22	Serum_Creatinine (mg/dL)	703 non-null	float64
23	Serum_Iron (μg/dL)	703 non-null	float64
24	Serum_Lactate (mmol/L)	703 non-null	float64
25	Serum_Urea (mg/dL)	703 non-null	float64
26	Serum_Zinc (μg/dL)	703 non-null	float64
27	TIBC (μg/dL)	703 non-null	float64
28	Transferrin_Saturation (%)	703 non-null	float64
29	WBC (10 ⁹ /L)	703 non-null	float64
30	pH	703 non-null	float64
31	Diagnosis	703 non-null	object

```
dtypes: float64(25), int64(2), object(5)
```

```
memory usage: 181.2+ KB
```

```
# Printing the first 10 rows of the dataframe with head(10)
```

```
# =====
```

```
df_train.head(10)
```

	index	AFP (ng/mL)	ALP (U/L)	ALT (U/L)	AST (U/L)	Age \
0	1124	13.571425	1653.138489	40.405592	45.598427	4
1	1188	13.649342	1940.518305	21.336986	34.064095	55
2	530	10.898943	1557.369920	29.665496	16.044488	30
3	686	13.872275	1273.840525	142.418649	64.204257	0
4	296	10.102457	1461.622515	22.437303	23.940205	59
5	663	5.824125	1921.018155	83.136847	39.372095	64
6	40	12.476249	2182.137388	97.396622	98.221842	29
7	889	11.157910	1804.807405	25.813858	59.671478	34
8	785	9.174636	1229.834674	27.789897	37.531030	3
9	1171	10.889718	1314.998314	33.851749	22.029857	50

	Albumin (g/dL)	Alcohol_Use (yes/no)	Bilirubin (mg/dL)	CRP (mg/L)
...	\			
0	4.477126	yes	0.541997	1.002121
...				
1	3.190724	yes	1.199063	0.582746
...				
2	4.506351	no	0.740952	1.670375
...				
3	3.665655	yes	1.939879	1.220646
...				
4	4.005109	no	0.769535	0.139863
...				
5	3.471696	no	0.557283	8.021424
...				
6	4.112928	no	1.670063	1.209236
...				
7	3.739201	yes	0.549854	0.980465
...				
8	4.359691	yes	0.749890	3.311809
...				
9	3.581977	yes	0.833413	1.065233
...				

	Serum_Creatinine (mg/dL)	Serum_Iron (µg/dL)	Serum_Lactate (mmol/L)
\			
0	2.569954	97.242495	2.923385
1	3.024642	128.597676	2.719966
2	2.219844	98.026938	3.267757
3	1.704282	63.105133	1.937625
4	2.656926	89.966713	1.979801

5	2.770285	96.542657	4.784693
6	2.801148	152.299112	1.687116
7	1.855184	72.934571	2.610542
8	3.953037	125.110690	2.343232
9	2.329074	40.899839	1.986068

	Serum_Urea (mg/dL)	Serum_Zinc (µg/dL)	TIBC (µg/dL) \
0	26.696110	96.124260	315.439318
1	23.021218	50.554650	268.492437
2	20.123808	82.900681	312.107620
3	36.436855	72.360010	182.506010
4	20.036074	63.216136	212.261519
5	27.760083	80.140859	304.210829
6	7.641927	70.282273	341.402911
7	16.227916	85.302620	201.170284
8	18.905484	81.647554	296.340915
9	20.626322	62.711051	339.683427

	Transferrin_Saturation (%)	WBC (10 ⁹ /L)	pH \
0	30.827639	4.180007	6.163600
1	47.896200	4.590995	5.244266
2	31.408057	4.524580	5.769592
3	34.577016	4.234483	5.982228
4	42.384844	4.306892	4.837404
5	31.735444	4.275489	5.187919
6	44.609787	8.069525	7.183443
7	36.255141	4.693278	7.304284
8	42.218500	4.526654	6.785159
9	12.040575	4.393664	5.718308

	Diagnosis
0	Healthy
1	Cirrhosis
2	Healthy
3	Drug-induced Liver Injury
4	Healthy
5	Fatty Liver Disease
6	Hepatitis
7	Fatty Liver Disease
8	Healthy
9	Healthy

[10 rows x 32 columns]

```
# Dropping the first column
# =====
df_train = df_train.drop('index', axis=1)

# Shwowing descriptive statistics of the dataset
# =====
df_train.describe()
```

	AFP (ng/mL)	ALP (U/L)	ALT (U/L)	AST (U/L)	Age \
count	703.000000	703.000000	703.000000	703.000000	703.000000
mean	23.804227	1624.337402	58.262199	46.855946	27.206259
std	54.938736	517.367426	36.372428	30.838523	23.632185
min	-22.836893	546.609117	-14.313734	-16.215535	0.000000
25%	8.029213	1229.735749	29.939330	24.163363	5.000000
50%	10.195074	1553.771408	48.857238	38.720602	23.000000
75%	12.656547	1931.433883	84.583900	70.023374	45.000000
max	449.925304	3769.332343	178.226798	154.173314	86.000000

	Albumin (g/dL)	Bilirubin (mg/dL)	CRP (mg/L)	Fibroscan (kPa)
count	703.000000	703.000000	703.000000	703.000000
mean	4.155568	1.056841	2.021190	8.899789
std	0.432317	0.617223	2.638333	3.822241
min	2.625845	0.130291	0.075035	2.323530
25%	3.880943	0.634089	0.651623	6.273365
50%	4.166995	0.801903	1.238123	7.966800
75%	4.437957	1.396853	2.449216	10.588808
max	5.447072	3.712815	32.356076	26.691648

	GGT (U/L)	...	Serum_Copper (µg/dL)	Serum_Creatinine (mg/dL) \
count	703.000000	...	703.000000	703.000000
mean	347.457426	...	105.754889	2.901900
std	371.157528	...	26.159403	0.647559
min	17.056568	...	30.099096	1.525736
25%	131.695741	...	88.632128	2.441051
50%	244.400101	...	102.842960	2.844621

```

75%      419.660403    ...      118.722674
3.275906
max      3426.981860    ...      227.617474
5.729147

```

	Serum_Iron (µg/dL)	Serum_Lactate (mmol/L)	Serum_Urea (mg/dL)
\			
count	703.000000	703.000000	703.000000
mean	105.901892	2.867523	21.440875
std	33.699054	0.936050	6.946169
min	8.328732	0.711997	7.641927
25%	84.737590	2.211484	16.422281
50%	104.670558	2.700933	20.668951
75%	125.937089	3.393898	25.096926
max	223.915730	8.010270	51.875943

	Serum_Zinc (µg/dL)	TIBC (µg/dL)	Transferrin_Saturation (%)	\
count	703.000000	703.000000	703.000000	
mean	78.021500	288.708468	38.220427	
std	15.826557	54.223207	15.438507	
min	34.170231	102.896996	2.496813	
25%	66.617374	248.832647	28.598082	
50%	78.024893	290.278094	35.617342	
75%	88.452369	326.508987	46.630605	
max	120.279580	484.581225	156.858874	

	WBC (10 ⁹ /L)	pH
count	703.000000	703.000000
mean	5.000254	6.072044
std	2.025527	0.642783
min	4.129884	4.416868
25%	4.243197	5.724293
50%	4.331605	6.085019
75%	4.762898	6.485114
max	21.195702	8.368522

[8 rows x 26 columns]

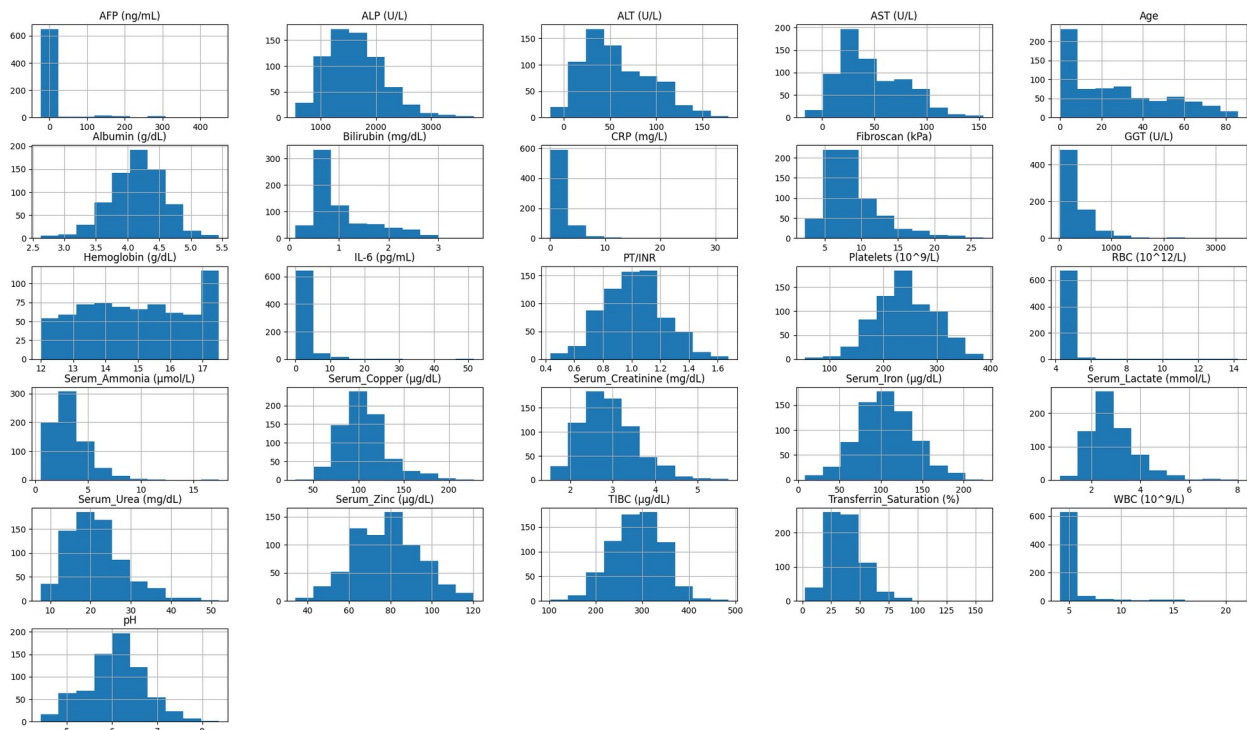
Histogram

```

# Plotting histograms for all columns in the dataset
# =====

```

```
df_train.hist(bins=10, figsize=(26, 15))
plt.show()
```

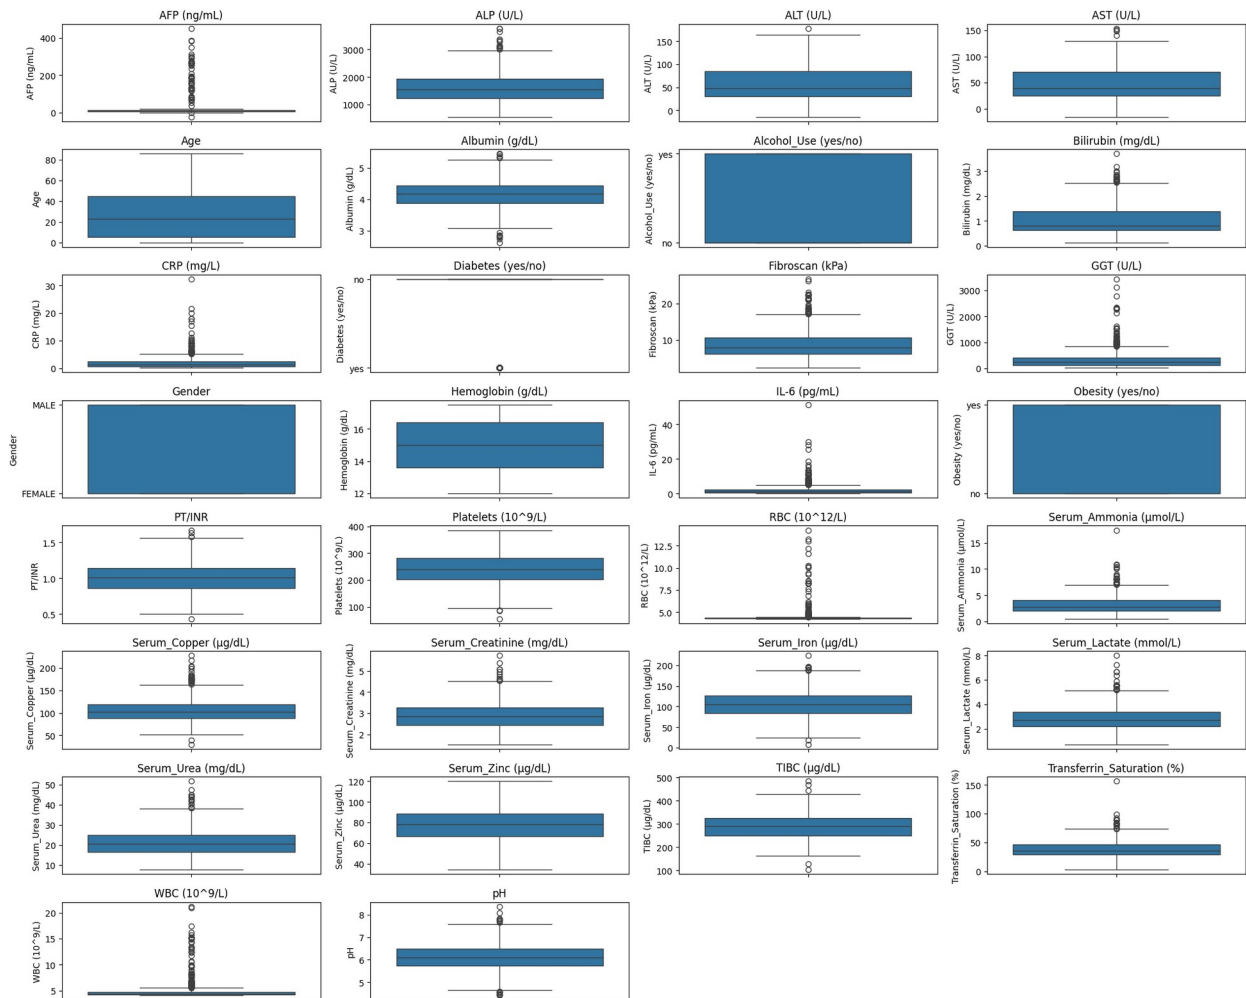


```
# Removing the target variable
# =====
features = df_train.columns[:-1]

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.boxplot(y=df_train[column])
    plt.title(column)
    plt.tight_layout()

plt.show()
```



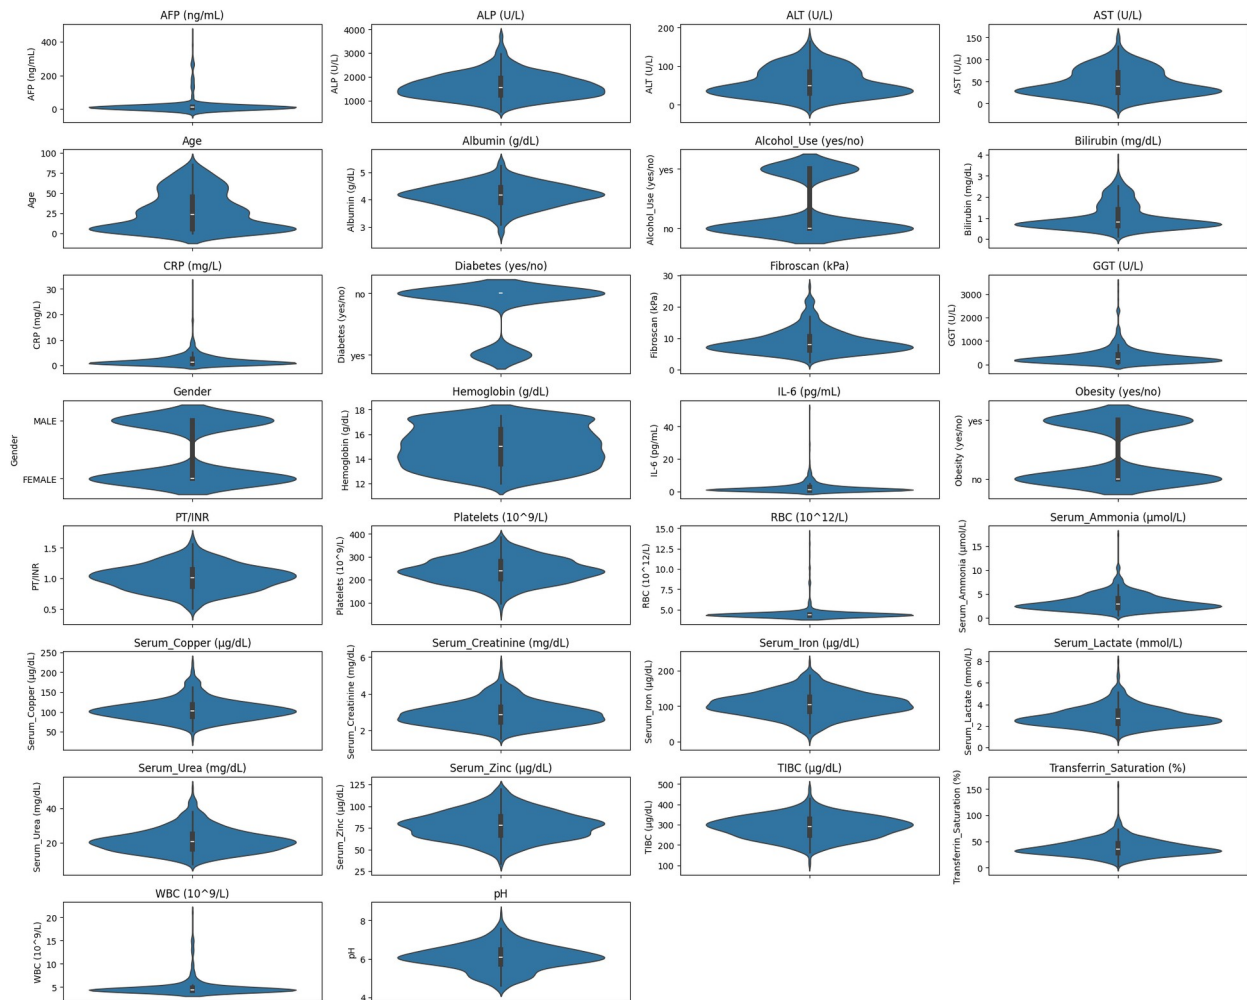
Violin Plots

```
# Removing the target variable
# =====
features = df_train.columns[:-1]

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.violinplot(y=df_train[column])
    plt.title(column)

plt.tight_layout()
plt.show()
```

Data cleaning

```
# Coping the dataset
```

```
# =====
```

```
df_train_copy = df_train.copy()
```

```
# printing the columns of the dataset
```

```
# =====
```

```
df_train_copy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 703 entries, 0 to 702
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	AFP (ng/mL)	703 non-null	float64
1	ALP (U/L)	703 non-null	float64
2	ALT (U/L)	703 non-null	float64
3	AST (U/L)	703 non-null	float64
4	Age	703 non-null	int64

5	Albumin (g/dL)	703	non-null	float64
6	Alcohol_Use (yes/no)	703	non-null	object
7	Bilirubin (mg/dL)	703	non-null	float64
8	CRP (mg/L)	703	non-null	float64
9	Diabetes (yes/no)	703	non-null	object
10	Fibroscan (kPa)	703	non-null	float64
11	GGT (U/L)	703	non-null	float64
12	Gender	703	non-null	object
13	Hemoglobin (g/dL)	703	non-null	float64
14	IL-6 (pg/mL)	703	non-null	float64
15	Obesity (yes/no)	703	non-null	object
16	PT/INR	703	non-null	float64
17	Platelets (10 ⁹ /L)	703	non-null	float64
18	RBC (10 ¹² /L)	703	non-null	float64
19	Serum_Ammonia (μmol/L)	703	non-null	float64
20	Serum_Copper (μg/dL)	703	non-null	float64
21	Serum_Creatinine (mg/dL)	703	non-null	float64
22	Serum_Iron (μg/dL)	703	non-null	float64
23	Serum_Lactate (mmol/L)	703	non-null	float64
24	Serum_Urea (mg/dL)	703	non-null	float64
25	Serum_Zinc (μg/dL)	703	non-null	float64
26	TIBC (μg/dL)	703	non-null	float64
27	Transferrin_Saturation (%)	703	non-null	float64
28	WBC (10 ⁹ /L)	703	non-null	float64
29	pH	703	non-null	float64
30	Diagnosis	703	non-null	object

dtypes: float64(25), int64(1), object(5)
memory usage: 175.8+ KB

Handling Catagorical Features

```
# Storing item of Diagnosis columns
# =====
diagnosis = df_train_copy['Diagnosis'].unique()

# Creating dictionary from array for diagnosis
# =====
diagnosis_dict = dict(zip(diagnosis, range(len(diagnosis))))

# Encoding the class label(Diagnosis) to integers
# =====
df_train_copy['Diagnosis'] =
df_train_copy['Diagnosis'].map(diagnosis_dict)

# Converting the catagorical data to numerical data
# =====
df_dummies = pd.get_dummies(df_train_copy, drop_first=True)
```

```

# Checking the info of the dataset
# =====
df_dummies.info()

<class 'pandas.core.frame.DataFrame'>
Index: 703 entries, 0 to 702
Data columns (total 31 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   AFP (ng/mL)                               703 non-null    float64
1   ALP (U/L)                                 703 non-null    float64
2   ALT (U/L)                                 703 non-null    float64
3   AST (U/L)                                 703 non-null    float64
4   Age                                         703 non-null    int64
5   Albumin (g/dL)                            703 non-null    float64
6   Bilirubin (mg/dL)                         703 non-null    float64
7   CRP (mg/L)                                703 non-null    float64
8   Fibroscan (kPa)                           703 non-null    float64
9   GGT (U/L)                                  703 non-null    float64
10  Hemoglobin (g/dL)                         703 non-null    float64
11  IL-6 (pg/mL)                              703 non-null    float64
12  PT/INR                                     703 non-null    float64
13  Platelets (10^9/L)                        703 non-null    float64
14  RBC (10^12/L)                             703 non-null    float64
15  Serum_Ammonia (μmol/L)                    703 non-null    float64
16  Serum_Copper (μg/dL)                      703 non-null    float64
17  Serum_Creatinine (mg/dL)                  703 non-null    float64
18  Serum_Iron (μg/dL)                        703 non-null    float64
19  Serum_Lactate (mmol/L)                    703 non-null    float64
20  Serum_Urea (mg/dL)                        703 non-null    float64
21  Serum_Zinc (μg/dL)                        703 non-null    float64
22  TIBC (μg/dL)                              703 non-null    float64
23  Transferrin_Saturation (%)                703 non-null    float64
24  WBC (10^9/L)                              703 non-null    float64
25  pH                                          703 non-null    float64
26  Diagnosis                                  703 non-null    int64
27  Alcohol_Use (yes/no)_yes                  703 non-null    bool
28  Diabetes (yes/no)_yes                     703 non-null    bool
29  Gender_MALE                               703 non-null    bool
30  Obesity (yes/no)_yes                      703 non-null    bool
dtypes: bool(4), float64(25), int64(2)
memory usage: 156.5 KB

# Moving the target variable to the last column
# =====
df_dummies = df_dummies[[col for col in df_dummies.columns if col !=
'Diagnosis']] + ['Diagnosis']
df_dummies.head()

```

AFP (ng/mL) (g/dL) \	ALP (U/L)	ALT (U/L)	AST (U/L)	Age	Albumin
0 13.571425	1653.138489	40.405592	45.598427	4	
4.477126					
1 13.649342	1940.518305	21.336986	34.064095	55	
3.190724					
2 10.898943	1557.369920	29.665496	16.044488	30	
4.506351					
3 13.872275	1273.840525	142.418649	64.204257	0	
3.665655					
4 10.102457	1461.622515	22.437303	23.940205	59	
4.005109					
Bilirubin (mg/dL)	CRP (mg/L)	Fibroscan (kPa)	GGT (U/L)	...	\
0 0.541997	1.002121	5.501881	311.253072	...	
1 1.199063	0.582746	13.902151	264.586392	...	
2 0.740952	1.670375	6.961641	66.247224	...	
3 1.939879	1.220646	9.193268	78.838414	...	
4 0.769535	0.139863	11.827885	246.940466	...	
Serum_Zinc (µg/dL)	TIBC (µg/dL)	Transferrin_Saturation (%)	WBC		
0 96.124260	315.439318		30.827639		
4.180007					
1 50.554650	268.492437		47.896200		
4.590995					
2 82.900681	312.107620		31.408057		
4.524580					
3 72.360010	182.506010		34.577016		
4.234483					
4 63.216136	212.261519		42.384844		
4.306892					
pH	Alcohol_Use (yes/no)_yes	Diabetes (yes/no)_yes			
Gender_MALE \					
0 6.163600	True	False			
True					
1 5.244266	True	False			
True					
2 5.769592	False	False			
False					
3 5.982228	True	False			
True					
4 4.837404	False	False			
False					
Obesity (yes/no)_yes	Diagnosis				
0 True	0				
1 True	1				
2 True	0				

3	False	2
4	True	0

[5 rows x 31 columns]

Removing Outliers

```
# Removing the outliers from the dataset using IQR method only for
numerical data
```

```
#
```

```
=====
```

```
# Taking only the numerical columns
```

```
df_numerical = df_dummies.select_dtypes(include=[np.number])
```

```
# Excluding the target variable
```

```
df_numerical = df_numerical.drop('Diagnosis', axis=1)
```

```
# Calculating the IQR
```

```
Q1 = df_numerical.quantile(0.25)
```

```
Q3 = df_numerical.quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
# Removing the outliers
```

```
df_numerical_2 = df_numerical[~((df_numerical < (Q1 - 1.5 * IQR)) |
(df_numerical > (Q3 + 1.5 * IQR))).any(axis=1)]
```

```
df_numerical_2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 309 entries, 0 to 701
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	AFP (ng/mL)	309 non-null	float64
1	ALP (U/L)	309 non-null	float64
2	ALT (U/L)	309 non-null	float64
3	AST (U/L)	309 non-null	float64
4	Age	309 non-null	int64
5	Albumin (g/dL)	309 non-null	float64
6	Bilirubin (mg/dL)	309 non-null	float64
7	CRP (mg/L)	309 non-null	float64
8	Fibroscan (kPa)	309 non-null	float64
9	GGT (U/L)	309 non-null	float64
10	Hemoglobin (g/dL)	309 non-null	float64
11	IL-6 (pg/mL)	309 non-null	float64
12	PT/INR	309 non-null	float64
13	Platelets (10^9/L)	309 non-null	float64
14	RBC (10^12/L)	309 non-null	float64
15	Serum_Ammonia (μmol/L)	309 non-null	float64
16	Serum_Copper (μg/dL)	309 non-null	float64
17	Serum_Creatinine (mg/dL)	309 non-null	float64
18	Serum_Iron (μg/dL)	309 non-null	float64
19	Serum_Lactate (mmol/L)	309 non-null	float64

```

20 Serum_Urea (mg/dL)          309 non-null    float64
21 Serum_Zinc (µg/dL)         309 non-null    float64
22 TIBC (µg/dL)               309 non-null    float64
23 Transferrin_Saturation (%)  309 non-null    float64
24 WBC (10^9/L)               309 non-null    float64
25 pH                         309 non-null    float64
dtypes: float64(25), int64(1)
memory usage: 65.2 KB

```

Using IQR method almost half of the data is getting eliminated so this method doesn't seem good one.

In the below I tried the z-score method to clean the data

```

# Detecting the outliers using the z-score (With the help of lecture
note)
#
=====
===
for column in df_numerical.columns:
    # Calculating the Z-scores for each column
    z_scores = (df_numerical[column] - df_numerical[column].mean()) /
df_numerical[column].std()

    # Detecting outliers using the absolute value of the Z-scores
(threshold of 2)
    outliers = (np.abs(z_scores) > 2)

    print(f"Number of outliers in {column}: {outliers.sum()}")

Number of outliers in AFP (ng/mL): 36
Number of outliers in ALP (U/L): 28
Number of outliers in ALT (U/L): 21
Number of outliers in AST (U/L): 24
Number of outliers in Age: 27
Number of outliers in Albumin (g/dL): 33
Number of outliers in Bilirubin (mg/dL): 45
Number of outliers in CRP (mg/L): 25
Number of outliers in Fibroscan (kPa): 40
Number of outliers in GGT (U/L): 26
Number of outliers in Hemoglobin (g/dL): 0
Number of outliers in IL-6 (pg/mL): 21
Number of outliers in PT/INR: 32
Number of outliers in Platelets (10^9/L): 28
Number of outliers in RBC (10^12/L): 19
Number of outliers in Serum_Ammonia (µmol/L): 28
Number of outliers in Serum_Copper (µg/dL): 35
Number of outliers in Serum_Creatinine (mg/dL): 33
Number of outliers in Serum_Iron (µg/dL): 36
Number of outliers in Serum_Lactate (mmol/L): 31

```

```
Number of outliers in Serum_Urea (mg/dL): 33
Number of outliers in Serum_Zinc (µg/dL): 29
Number of outliers in TIBC (µg/dL): 27
Number of outliers in Transferrin_Saturation (%): 29
Number of outliers in WBC (10^9/L): 29
Number of outliers in pH: 35
```

```
# Printing information of the dataset with numerical columns
```

```
# =====
df_numerical.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 703 entries, 0 to 702
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	AFP (ng/mL)	703 non-null	float64
1	ALP (U/L)	703 non-null	float64
2	ALT (U/L)	703 non-null	float64
3	AST (U/L)	703 non-null	float64
4	Age	703 non-null	int64
5	Albumin (g/dL)	703 non-null	float64
6	Bilirubin (mg/dL)	703 non-null	float64
7	CRP (mg/L)	703 non-null	float64
8	Fibroscan (kPa)	703 non-null	float64
9	GGT (U/L)	703 non-null	float64
10	Hemoglobin (g/dL)	703 non-null	float64
11	IL-6 (pg/mL)	703 non-null	float64
12	PT/INR	703 non-null	float64
13	Platelets (10^9/L)	703 non-null	float64
14	RBC (10^12/L)	703 non-null	float64
15	Serum_Ammonia (µmol/L)	703 non-null	float64
16	Serum_Copper (µg/dL)	703 non-null	float64
17	Serum_Creatinine (mg/dL)	703 non-null	float64
18	Serum_Iron (µg/dL)	703 non-null	float64
19	Serum_Lactate (mmol/L)	703 non-null	float64
20	Serum_Urea (mg/dL)	703 non-null	float64
21	Serum_Zinc (µg/dL)	703 non-null	float64
22	TIBC (µg/dL)	703 non-null	float64
23	Transferrin_Saturation (%)	703 non-null	float64
24	WBC (10^9/L)	703 non-null	float64
25	pH	703 non-null	float64

```
dtypes: float64(25), int64(1)
```

```
memory usage: 148.3 KB
```

```
# Removing the outliers from the dataset using Z-score method only for
numerical data (With the help of lecture note)
```

```
#
```

```
=====
=====
```

```

for column in df_numerical.columns:
    # Calculate the z-scores for each column
    z_scores = (df_numerical[column] - df_numerical[column].mean()) /
df_numerical[column].std()

    # Only keep rows in dataframe where the z-score is less than 2
    standard deviations
    df_numerical_3 = df_numerical[np.abs(z_scores) < 2]

# Checking the information of the dataset
df_numerical_3.info()

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 668 entries, 0 to 702
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	AFP (ng/mL)	668 non-null	float64
1	ALP (U/L)	668 non-null	float64
2	ALT (U/L)	668 non-null	float64
3	AST (U/L)	668 non-null	float64
4	Age	668 non-null	int64
5	Albumin (g/dL)	668 non-null	float64
6	Bilirubin (mg/dL)	668 non-null	float64
7	CRP (mg/L)	668 non-null	float64
8	Fibroscan (kPa)	668 non-null	float64
9	GGT (U/L)	668 non-null	float64
10	Hemoglobin (g/dL)	668 non-null	float64
11	IL-6 (pg/mL)	668 non-null	float64
12	PT/INR	668 non-null	float64
13	Platelets (10 ⁹ /L)	668 non-null	float64
14	RBC (10 ¹² /L)	668 non-null	float64
15	Serum_Ammonia (μmol/L)	668 non-null	float64
16	Serum_Copper (μg/dL)	668 non-null	float64
17	Serum_Creatinine (mg/dL)	668 non-null	float64
18	Serum_Iron (μg/dL)	668 non-null	float64
19	Serum_Lactate (mmol/L)	668 non-null	float64
20	Serum_Urea (mg/dL)	668 non-null	float64
21	Serum_Zinc (μg/dL)	668 non-null	float64
22	TIBC (μg/dL)	668 non-null	float64
23	Transferrin_Saturation (%)	668 non-null	float64
24	WBC (10 ⁹ /L)	668 non-null	float64
25	pH	668 non-null	float64

```
dtypes: float64(25), int64(1)
```

```
memory usage: 140.9 KB
```



```
dtypes: bool(4), float64(25), int64(2)
memory usage: 148.7 KB
```

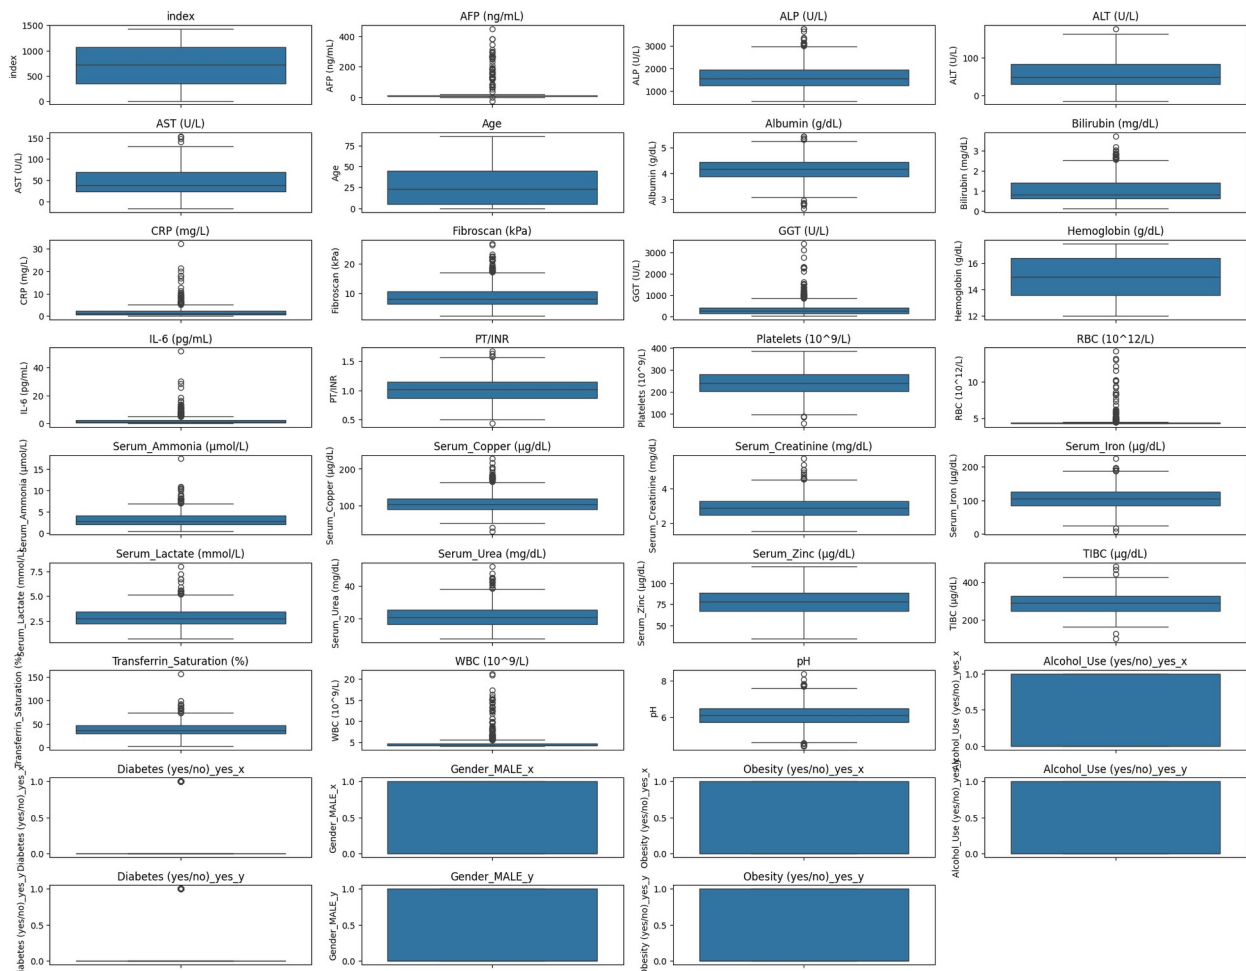
Visualizing after cleaning

```
# Removing the target variable
# =====
features = df_cleaned.columns[:-1]

plt.figure(figsize=(20, 60))

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.boxplot(y=df_cleaned[column])
    plt.title(column)
    plt.tight_layout()

plt.show()
```



Still shows a lot of Outliers.

Data preprocessing and visualisation

Splitting the Dataset

```
#-----  
# Splitting the cleaned dataset into test and train  
#-----  
X_cleaned = df_cleaned.drop('Diagnosis', axis=1)  
y_cleaned = df_cleaned['Diagnosis']  
  
X_train_cleaned, X_test_cleaned, y_train_cleaned, y_test_cleaned =  
train_test_split(X_cleaned, y_cleaned, test_size=0.2, random_state=42)  
  
#-----  
# Splitting the dataset with outlier into test and train  
#-----  
X = df_dummies.drop('Diagnosis', axis=1)  
y = df_dummies['Diagnosis']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

Exploring training data after cleaning and splitting

```
X_cleaned.info()  
  
<class 'pandas.core.frame.DataFrame'>  
Index: 668 entries, 0 to 702  
Data columns (total 30 columns):  
#   Column                                     Non-Null Count  Dtype  
---  -  
0   AFP (ng/mL)                               668 non-null    float64  
1   ALP (U/L)                                 668 non-null    float64  
2   ALT (U/L)                                 668 non-null    float64  
3   AST (U/L)                                 668 non-null    float64  
4   Age                                         668 non-null    int64  
5   Albumin (g/dL)                            668 non-null    float64  
6   Bilirubin (mg/dL)                         668 non-null    float64  
7   CRP (mg/L)                                668 non-null    float64  
8   Fibroscan (kPa)                           668 non-null    float64  
9   GGT (U/L)                                 668 non-null    float64  
10  Hemoglobin (g/dL)                         668 non-null    float64  
11  IL-6 (pg/mL)                              668 non-null    float64  
12  PT/INR                                     668 non-null    float64  
13  Platelets (10^9/L)                        668 non-null    float64  
14  RBC (10^12/L)                             668 non-null    float64  
15  Serum_Ammonia (µmol/L)                    668 non-null    float64  
16  Serum_Copper (µg/dL)                      668 non-null    float64  
17  Serum_Creatinine (mg/dL)                  668 non-null    float64
```

```

18 Serum_Iron (µg/dL)          668 non-null    float64
19 Serum_Lactate (mmol/L)      668 non-null    float64
20 Serum_Urea (mg/dL)          668 non-null    float64
21 Serum_Zinc (µg/dL)          668 non-null    float64
22 TIBC (µg/dL)                668 non-null    float64
23 Transferrin_Saturation (%)  668 non-null    float64
24 WBC (10^9/L)                668 non-null    float64
25 pH                          668 non-null    float64
26 Alcohol_Use (yes/no)_yes    668 non-null    bool
27 Diabetes (yes/no)_yes       668 non-null    bool
28 Gender_MALE                  668 non-null    bool
29 Obesity (yes/no)_yes        668 non-null    bool
dtypes: bool(4), float64(25), int64(1)
memory usage: 143.5 KB

```

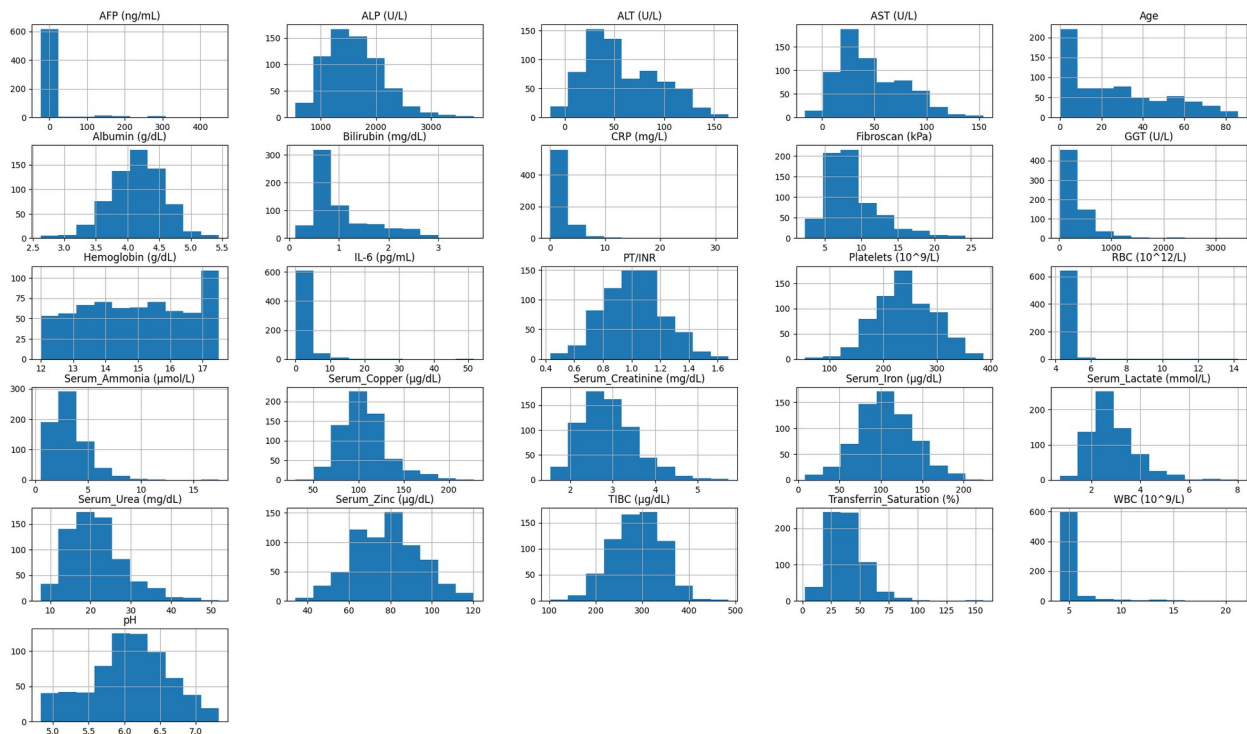
```
# Plotting histograms for all columns in the dataset
```

```
# =====
```

```

X_cleaned.hist(bins=10, figsize=(26, 15))
plt.show()

```



```
# Plotting the Boxplot for the cleaned dataset i.e
```

```
X_train_cleaned_sc_df
```

```
#
```

```
=====
```

```
features = X_cleaned.columns[:-1]
```

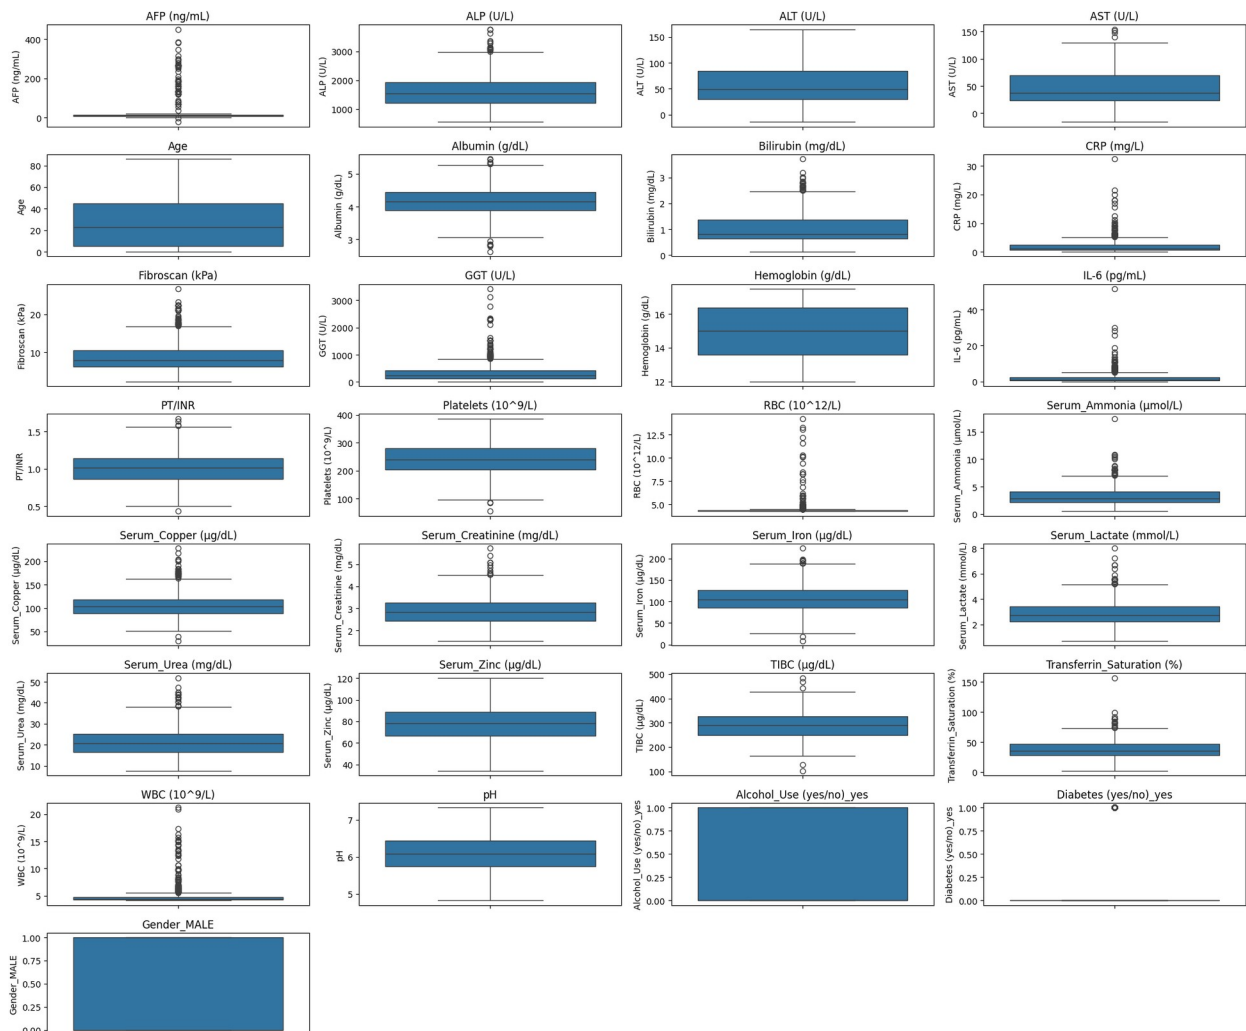
```
plt.figure(figsize=(20, 60))
```

```

for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.boxplot(y=X_cleaned[column])
    plt.title(column)
    plt.tight_layout()

plt.show()

```



```

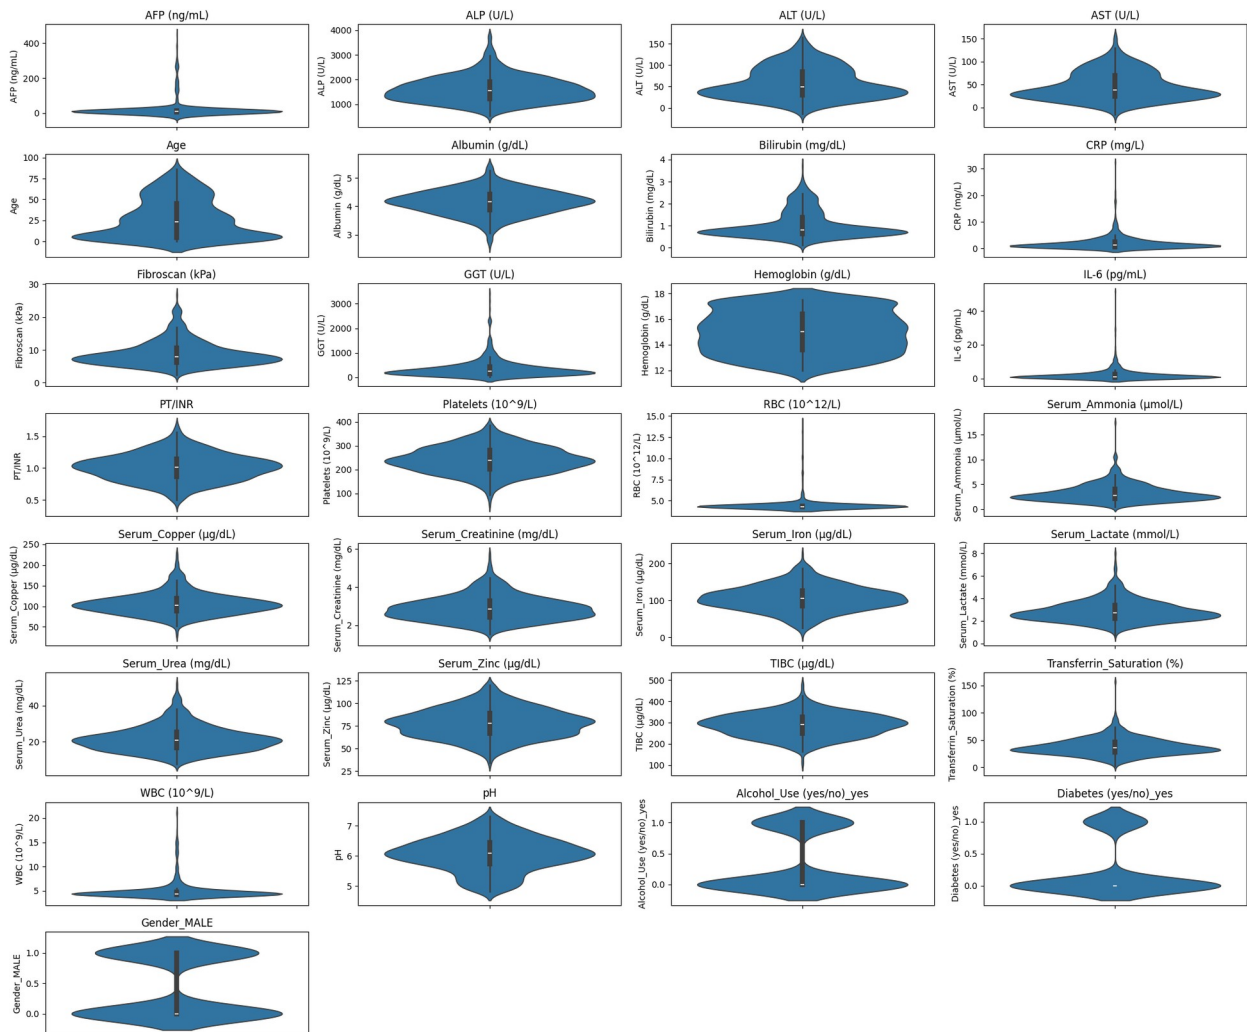
# Plotting the Violin plot for the cleaned dataset i.e
X_train_cleaned_sc_df
#
=====
====
features = X_cleaned.columns[:-1]

plt.figure(figsize=(20, 60))

```

```
for i, column in enumerate(features):
    plt.subplot(len(features), 4, i + 1)
    sns.violinplot(y=X_cleaned[column])
    plt.title(column)
```

```
plt.tight_layout()
plt.show()
```



Modelling

Pipelining with Kernel

```
# Defining the pipeline with kernel(SVC) and pca
# =====
pipe_svc_pca = Pipeline([
    ('scaler', StandardScaler()),
    ('pca', PCA()),
    ('classifier', SVC())
])
```

```

])

# Defining the parameters for the grid search
# =====
param_grid_svc_pca = {
    'pca__n_components': [0.9, 0.95, 0.99],
    'classifier__C': [0.1, 1, 10],
    'classifier__kernel': ['linear', 'poly', 'rbf', 'sigmoid']
}

# Defining the grid search
# =====
grid_search_svc_pca = GridSearchCV(estimator=pipe_svc_pca,
    param_grid=param_grid_svc_pca, scoring='f1_macro', cv=5)

# Fitting the grid search
# =====
grid_search_svc_pca.fit(X_train_cleaned, y_train_cleaned)

GridSearchCV(cv=5,
    estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                ('pca', PCA()), ('classifier',
SVC())]),
    param_grid={'classifier__C': [0.1, 1, 10],
                'classifier__kernel': ['linear', 'poly',
'rbf',
                                'sigmoid'],
                'pca__n_components': [0.9, 0.95, 0.99]},
    scoring='f1_macro')

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_svc_pca.best_params_}")

# Printing the best score
# =====
print(f"Best score: {grid_search_svc_pca.best_score_}")

Best parameters: {'classifier__C': 1, 'classifier__kernel': 'sigmoid',
'pca__n_components': 0.95}
Best score: 0.7424156305547781

# Defining the pipeline with kernel(SVC) and lda
# =====
pipe_svc_lda = Pipeline([
    ('scaler', StandardScaler()),
    ('lda', LDA()),
    ('classifier', SVC())
])

# Defining the parameters for the grid search

```

```

# =====
param_grid_svc_lda= {
    'lda__n_components': [1, 2],
    'classifier__C': [0.1, 1, 10],
    'classifier__kernel': ['linear', 'poly', 'rbf', 'sigmoid']
}

# Defining the grid search
# =====
grid_search_svc_lda = GridSearchCV(estimator=pipe_svc_lda,
    param_grid=param_grid_svc_lda, scoring='f1_macro', cv=5)

# Fitting the grid search
# =====
grid_search_svc_lda.fit(X_train_cleaned, y_train_cleaned)

GridSearchCV(cv=5,
    estimator=Pipeline(steps=[('scaler', StandardScaler()),
    ('lda',
    LinearDiscriminantAnalysis()),
    ('classifier', SVC())]),
    param_grid={'classifier__C': [0.1, 1, 10],
    'classifier__kernel': ['linear', 'poly',
    'rbf',
    'sigmoid'],
    'lda__n_components': [1, 2]},
    scoring='f1_macro')

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_svc_lda.best_params_}")

# Printing the best score
# =====
print(f"Best score: {grid_search_svc_lda.best_score_}")

Best parameters: {'classifier__C': 10, 'classifier__kernel': 'linear',
'lda__n_components': 2}
Best score: 0.4884996676523051

```

Pipelining with Regularization

```

# Defining the pipeline with Regularization(LogisticRegression) and pca
# =====
pipe_lr_pca = Pipeline([
    ('scaler', StandardScaler()),
    ('pca', PCA()),
    ('classifier', LogisticRegression())
])

# Defining the parameters for the grid search

```



```

# =====
param_grid_lr_pca = {
    'pca_n_components': [0.9, 0.95, 0.99],
    'classifier_C': [0.1, 1, 10],
    'classifier_penalty': ['l1', 'l2']
}

# Defining the grid search
# =====
grid_search_lr_pca = GridSearchCV(estimator=pipe_lr_pca,
    param_grid=param_grid_lr_pca, scoring='f1_macro', cv=5)

# Fitting the grid search
# =====
grid_search_lr_pca.fit(X_train_cleaned, y_train_cleaned)

/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_validation.py:542: FitFailedWarning:
45 fits failed out of a total of 90.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

Below are more details about the failures:
-----
-----
45 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_validation.py", line 890, in
_fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/base.py", line 1351, in wrapper
    return fit_method(estimator, *args, **kwargs)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/pipeline.py", line 475, in fit
    self._final_estimator.fit(Xt, y, **last_step_params["fit"])
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/base.py", line 1351, in wrapper
    return fit_method(estimator, *args, **kwargs)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py", line 1172, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py", line 67, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1

```

penalty.

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_search.py:1051: UserWarning: One or
more of the test scores are non-finite: [          nan          nan
nan 0.73142056 0.73013276 0.74056691
          nan          nan          nan 0.70825524 0.70806467 0.71615276
          nan          nan          nan 0.68664242 0.6847317  0.69363941]
warnings.warn(
```

```
GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                       ('pca', PCA()),
                                       ('classifier',
LogisticRegression())]),
             param_grid={'classifier__C': [0.1, 1, 10],
                        'classifier__penalty': ['l1', 'l2'],
                        'pca__n_components': [0.9, 0.95, 0.99]},
             scoring='f1_macro')
```

```
# Printing the best parameters
```

```
# =====
```

```
print(f"Best parameters: {grid_search_lr_pca.best_params_}")
```

```
# Printing the best score
```

```
# =====
```

```
print(f"Best score: {grid_search_lr_pca.best_score_}")
```

```
Best parameters: {'classifier__C': 0.1, 'classifier__penalty': 'l2',
'pca__n_components': 0.99}
```

```
Best score: 0.7405669117166076
```

```
# Defining the pipeline with Regularization(LogisticRegression) and lda
```

```
# =====
```

```
pipe_lr_lda = Pipeline([
    ('scaler', StandardScaler()),
    ('lda', LDA()),
    ('classifier', LogisticRegression())
])
```

```
# Defining the parameters for the grid search
```

```
# =====
```

```
param_grid_lr_lda = {
    'lda__n_components': [1, 2],
    'classifier__C': [0.001, 0.01, 0.1, 1, 10, 100],
    'classifier__penalty': ['l1', 'l2']
}
```

```
# Defining the grid search
```

```
# =====
```

```

grid_search_lr_lda = GridSearchCV(estimator=pipe_lr_lda,
param_grid=param_grid_lr_lda, scoring='f1_macro', cv=5)

# Fitting the grid search
# =====
grid_search_lr_lda.fit(X_train_cleaned, y_train_cleaned)

/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_validation.py:542: FitFailedWarning:
60 fits failed out of a total of 120.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.

Below are more details about the failures:
-----
-----
60 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_validation.py", line 890, in
_fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/base.py", line 1351, in wrapper
    return fit_method(estimator, *args, **kwargs)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/pipeline.py", line 475, in fit
    self._final_estimator.fit(Xt, y, **last_step_params["fit"])
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/base.py", line 1351, in wrapper
    return fit_method(estimator, *args, **kwargs)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py", line 1172, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py", line 67, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1
penalty.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/model_selection/_search.py:1051: UserWarning: One or
more of the test scores are non-finite: [      nan      nan
0.20646821 0.22460156      nan      nan
0.25982519 0.37231984      nan      nan 0.29688083 0.47922055
      nan      nan 0.30158403 0.49746643      nan      nan

```

```

0.31018654 0.50354344 nan nan 0.30933998 0.50567583]
warnings.warn(

GridSearchCV(cv=5,
              estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                         ('lda',
LinearDiscriminantAnalysis()),
                                         ('classifier',
LogisticRegression())]),
              param_grid={'classifier__C': [0.001, 0.01, 0.1, 1, 10,
100],
                          'classifier__penalty': ['l1', 'l2'],
                          'lda__n_components': [1, 2]},
              scoring='f1_macro')

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_lr_lda.best_params}")

# Printing the best score
# =====
print(f"Best score: {grid_search_lr_lda.best_score}")

Best parameters: {'classifier__C': 100, 'classifier__penalty': 'l2',
'lda__n_components': 2}
Best score: 0.50567582960442

# Defining the pipeline with RandomForestClassifier and pca
# =====
pipe_rf_pca = Pipeline([
    ('scaler', StandardScaler()),
    ('pca', PCA()),
    ('classifier', RandomForestClassifier())
])

# Defining the parameters for the grid search with f1 macro average as
scoring
#
=====
=====
param_grid_rf_pca = {
    'pca__n_components': [0.9, 0.95, 0.99],
    'classifier__n_estimators': [100, 200, 300],
    'classifier__max_depth': [5, 10, 15, 20],
}

# Defining the grid search
# =====
grid_search_rf_pca = GridSearchCV(estimator=pipe_rf_pca,
param_grid=param_grid_rf_pca, scoring='f1_macro', cv=5)

```

```

# Fitting the grid search(for cleaned dataset)
# =====
grid_search_rf_pca.fit(X_train_cleaned, y_train_cleaned)

GridSearchCV(cv=5,
              estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                         ('pca', PCA()),
                                         ('classifier',
                                          RandomForestClassifier())]),
              param_grid={'classifier__max_depth': [5, 10, 15, 20],
                          'classifier__n_estimators': [100, 200, 300],
                          'pca__n_components': [0.9, 0.95, 0.99]},
              scoring='f1_macro')

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_rf_pca.best_params}")

# Printing the best score
# =====
print(f"Best score: {grid_search_rf_pca.best_score}")

Best parameters: {'classifier__max_depth': 20,
                  'classifier__n_estimators': 300, 'pca__n_components': 0.99}
Best score: 0.6654014631167522

# Defining the pipeline with RandomForestClassifier and lda
# =====
pipe_rf_lda = Pipeline([
    ('scaler', StandardScaler()),
    ('lda', LDA()),
    ('classifier', RandomForestClassifier())
])

# Defining the parameters for the grid search
# =====
param_grid_rf_lda = {
    'lda__n_components': [1, 2],
    'classifier__n_estimators': [100, 200, 300],
    'classifier__max_depth': [5, 10, 15, 20]
}

# Defining the grid search
# =====
grid_search_rf_lda = GridSearchCV(estimator=pipe_rf_lda,
                                   param_grid=param_grid_rf_lda, scoring='f1_macro', cv=5)

# Fitting the grid search(for cleaned dataset)
# =====
grid_search_rf_lda.fit(X_train_cleaned, y_train_cleaned)

```

```

GridSearchCV(cv=5,
             estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                       ('lda',
                                        LinearDiscriminantAnalysis()),
                                       ('classifier',
                                        RandomForestClassifier())]),
             param_grid={'classifier__max_depth': [5, 10, 15, 20],
                         'classifier__n_estimators': [100, 200, 300],
                         'lda__n_components': [1, 2]},
             scoring='f1_macro')

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_rf_lda.best_params}")

# Printing the best score
# =====
print(f"Best score: {grid_search_rf_lda.best_score}")

Best parameters: {'classifier__max_depth': 10,
                  'classifier__n_estimators': 200, 'lda__n_components': 2}
Best score: 0.5107671041902294

# Defining the pipeline with RandomForestClassifier
# =====
pipe_rf = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', RandomForestClassifier())
])

# Defining the parameters for the grid search
# =====
param_grid_rf = {
    'classifier__n_estimators': [100, 200, 300],
    'classifier__max_depth': [5, 10, 15, 20]
}

# Defining the grid search
# =====
grid_search_rf = GridSearchCV(estimator=pipe_rf,
                              param_grid=param_grid_rf, scoring='f1_macro', cv=5)

# Fitting the grid search(for cleaned dataset)
# =====
grid_search_rf.fit(X_train_cleaned, y_train_cleaned)

# Printing the best parameters
# =====
print(f"Best parameters: {grid_search_rf.best_params}")

# Printing the best score

```

```
# =====  
print(f"Best score: {grid_search_rf.best_score}")  
  
Best parameters: {'classifier__max_depth': 15,  
                  'classifier__n_estimators': 300}  
Best score: 0.7469526003955284
```

Confusion Matrix for the best model with 60-40 split

```
# Predicting the test data best model  
# =====  
X_train_c, X_test_c, y_train_c, y_test_c = train_test_split(X, y,  
                                                            test_size=0.2, random_state=42)  
best_model = grid_search_svc_pca.best_estimator_  
best_model.fit(X_train_c, y_train_c)  
y_pred = best_model.predict(X_test_c)  
  
# Generating the confusion matrix (Taken from the lecture)  
# =====  
ConfusionMatrixDisplay.from_predictions(y_test_c, y_pred)  
plt.title('Confusion Matrix for SVC with PCA')  
plt.xticks(rotation=90)  
plt.show()
```


0	AFP (ng/mL)	668 non-null	float64
1	ALP (U/L)	668 non-null	float64
2	ALT (U/L)	668 non-null	float64
3	AST (U/L)	668 non-null	float64
4	Age	668 non-null	int64
5	Albumin (g/dL)	668 non-null	float64
6	Bilirubin (mg/dL)	668 non-null	float64
7	CRP (mg/L)	668 non-null	float64
8	Fibroscan (kPa)	668 non-null	float64
9	GGT (U/L)	668 non-null	float64
10	Hemoglobin (g/dL)	668 non-null	float64
11	IL-6 (pg/mL)	668 non-null	float64
12	PT/INR	668 non-null	float64
13	Platelets (10^9/L)	668 non-null	float64
14	RBC (10^12/L)	668 non-null	float64
15	Serum_Ammonia (μmol/L)	668 non-null	float64
16	Serum_Copper (μg/dL)	668 non-null	float64
17	Serum_Creatinine (mg/dL)	668 non-null	float64
18	Serum_Iron (μg/dL)	668 non-null	float64
19	Serum_Lactate (mmol/L)	668 non-null	float64
20	Serum_Urea (mg/dL)	668 non-null	float64
21	Serum_Zinc (μg/dL)	668 non-null	float64
22	TIBC (μg/dL)	668 non-null	float64
23	Transferrin_Saturation (%)	668 non-null	float64
24	WBC (10^9/L)	668 non-null	float64
25	pH	668 non-null	float64
26	Alcohol_Use (yes/no)_yes	668 non-null	bool
27	Diabetes (yes/no)_yes	668 non-null	bool
28	Gender_MALE	668 non-null	bool
29	Obesity (yes/no)_yes	668 non-null	bool

dtypes: bool(4), float64(25), int64(1)

memory usage: 143.5 KB

df_test_dummies.info()

<class 'pandas.core.frame.DataFrame'>

Index: 302 entries, 0 to 301

Data columns (total 30 columns):

#	Column	Non-Null Count	Dtype
0	AFP (ng/mL)	302 non-null	float64
1	ALP (U/L)	302 non-null	float64
2	ALT (U/L)	302 non-null	float64
3	AST (U/L)	302 non-null	float64
4	Age	302 non-null	int64
5	Albumin (g/dL)	302 non-null	float64
6	Bilirubin (mg/dL)	302 non-null	float64
7	CRP (mg/L)	302 non-null	float64
8	Fibroscan (kPa)	302 non-null	float64

9	GGT (U/L)	302	non-null	float64
10	Hemoglobin (g/dL)	302	non-null	float64
11	IL-6 (pg/mL)	302	non-null	float64
12	PT/INR	302	non-null	float64
13	Platelets (10 ⁹ /L)	302	non-null	float64
14	RBC (10 ¹² /L)	302	non-null	float64
15	Serum_Ammonia (μmol/L)	302	non-null	float64
16	Serum_Copper (μg/dL)	302	non-null	float64
17	Serum_Creatinine (mg/dL)	302	non-null	float64
18	Serum_Iron (μg/dL)	302	non-null	float64
19	Serum_Lactate (mmol/L)	302	non-null	float64
20	Serum_Urea (mg/dL)	302	non-null	float64
21	Serum_Zinc (μg/dL)	302	non-null	float64
22	TIBC (μg/dL)	302	non-null	float64
23	Transferrin_Saturation (%)	302	non-null	float64
24	WBC (10 ⁹ /L)	302	non-null	float64
25	pH	302	non-null	float64
26	Alcohol_Use (yes/no)_yes	302	non-null	bool
27	Diabetes (yes/no)_yes	302	non-null	bool
28	Gender_MALE	302	non-null	bool
29	Obesity (yes/no)_yes	302	non-null	bool

dtypes: bool(4), float64(25), int64(1)
memory usage: 64.9 KB

```
# Fitting the best model on the training data
# =====
best_model.fit(X_cleaned, y_cleaned)

# Predicting the test data
# =====
y_pred_final = best_model.predict(df_test_dummies)
```

Kaggle submission

```
# Mapping the diagnosis back to the original values
# =====
diagonosis_dict_reverse = {v: k for k, v in diagnosis_dict.items()}
y_pred_final = pd.Series(y_pred_final).map(diagonosis_dict_reverse)

#-----
# Creating csv file for Kaggle submission
#-----
DF = pd.DataFrame(y_pred_final)
headers = ["Diagnosis"]
DF.columns = headers
DF.to_csv('predictions.csv', index_label='index', index=True)
```

ROC Curve for Binary Version of the Target variable

```
# Converting the diagnosis column to 2 classes
# =====
df_train_simplified = df_cleaned.copy()
df_train_simplified['Diagnosis'] =
df_train_simplified['Diagnosis'].apply(lambda x: 0 if x == 0 else 1)

X_simplified = df_train_simplified.drop('Diagnosis', axis=1)
y_simplified = df_train_simplified['Diagnosis']

# Defining the classifier
# =====
lr = LogisticRegression()
# Producing the probabilities for each class
y_prob = cross_val_predict(lr, X_simplified, y_simplified, cv=5,
method='predict_proba')
fpr, tpr, thresholds = roc_curve(y_simplified, y_prob[:, 1])

# Plotting the ROC curve
# =====
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.show()
```

```
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```
regression
  n_iter_i = _check_optimize_result(
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
  n_iter_i = _check_optimize_result(
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
  n_iter_i = _check_optimize_result(
/Users/jony/opt/anaconda3/envs/dat200_env/lib/python3.10/site-
packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
  n_iter_i = _check_optimize_result(
```

